

HWX Ticket Details

ReD Shield Engineering

Exported on 06/05/2020

Table of Contents


HWX Patch details:

Date	Patch #	Patch Details	Remarks
 29 Jan 2019	3.1.0.0-78	HDP 2.6.5 to 3.1	upgraded to HDP 3.1 from 2.6.5
 11 Mar 2019	3.1.0.10-1	HiveWarehouse connector not able to write data into complex data types in Hive acid tables In Spark Structured Streaming Column containing timestamp value is not stored with proper values in hive table with column type timestamp or string	HWX Case #00207534 HWX Case #00207953 ¹
 29 Apr 2019	3.1.0.22-3	Spark jobs failing with Memory leak issue in Hive Join results in ArrayIndexOutOfBoundsException , when used executeQuery()	HWX Case #216491 HWX Case #212771
 18 Jul 2019	3.1.0.103-1	vertex issues	HWX Case #213668 ²
 09 Aug 2019	3.1.0.126-1	Delta folder not found and transformation job fails. Issue with the compaction's race condition.	HWX Case #00224459 ³
 28 Aug 2019	3.1.0.139-1	LLAP Containers unstable - Fixed through dev jar #1 Compaction race condition's patch kills the ingestion job which is fixed through the dev jar#2. This official patch fixes the both dev jars combined together	HWX Case # 00234501
 20 Sep 2019	3.1.0.158-1	This is the official patch for the batch transformation jobs slowness	HWX Case # 00234501
 21 Oct 2019	3.1.0.175-5	AOD clusters reported Caused by: java.lang.ArrayIndexOutOfBoundsException: also multiple errors on their spark submit jobs.	AOD HWX Cases 00242128 00236938 00238255.


¹ <https://support.hortonworks.com/s/case/500440000108bHnAAI/detail>

² <https://support.hortonworks.com/s/case/500440000120jM3AAK/detail>




³ <https://support.hortonworks.com/s/case/5002R000013NFoYQAW/detail>





Date	Patch #	Patch Details	Remarks
 17 Feb 2020	DEV Jar Patch	Vectorization issues reported for IBI performance issues. Received the DEV jar from Cloudera for the same md5sum of DEV Jar patch 55e740de4167e3a8a7d71a515da5f11e → hive-exec-3.1.0.3.1.0.175-5.jar	HWX Case # 653597







This page contains HWX Ticket details:



S . No	Ticket #	Problem Description	Date Opened	Current Status	Remarks	Component	Type of Ticket	Ticket Created By	Ticket Attended By	Comments	Remarks
186	669089	Not able to initiate compaction of segments on druid datasource	 20 Mar 2020	Closed	Resolved	Druid	S2 - Core Feature Inoperative	Murugesan Kumar	Nisha	Not able to initiate compaction of segments on druid datasource. Cloudera provided the solution: Ref: [1] https://druid.apache.org/docs/latest/tutorials/tutorial-compaction.html [2] https://druid.apache.org/docs/0.15.1-incubating/ingestion/compaction.html ⁴	fixed

⁴ <https://druid.apache.org/docs/0.15.1-incubating/ingestion/compaction.html>




185	667695	Unexpected reboot crashed the Ambari with Postgres Database	 15 Mar 2020	Closed	Closed	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Anchal	<p>Coventry data centre nodes are getting restarted unexpectedly. Due that we are not able to bring up the postgres services with Ambari service.</p> <p>Corp-IT team fixed the issue.</p> <p>It shows <code>/var/run/postgresql*</code> is missing there, need to create it and change the ownership to postgres.</p> <p>Now started Ambari and other Hadoop services too</p>	fixed by Corp-IT team
184	666380	Kafka indexing service Druid not working with Secured kafka with Kerberos	 10 Mar 2020	Closed	Closed	Druid	S2 - Core Feature Inoperative	Murugesan Kumar	Mariano	<p>We need to start ingestion directly from kafka(Kerberos enabled and SASL_SSL connection) to Druid data source. We have created the supervisor spec and also have set the java opt parameters.</p> <p>Unable to access the Druid, shared the json file in the case too. Waiting for CLDR support engineer to be assigned.</p> <p>Mariano fixed that and now able to store the data in the data source .</p>	
183	664063	Druid can not read from Kafka in Dev	 28 Feb 2020	Closed	Closed	Druid - >Kafka	S2 - Core Feature Inoperative	Phil Young	Phil Young	<p>While attempting to induce Druid to read from kafka, we got an error that suggests that support for Kafka is not enabled in Druid (in the dev cluster).</p> <p>CLDR Support provided the following ref and closed the ticket: Ref: https://druid.apache.org/docs/latest/development/extensions.html#loading-extensions https://druid.apache.org/docs/latest/development/extensions-core/kafka-ingestion.html</p>	


182	658797	LLAP crashed for weekend run and unable to find the Root Cause	 23 Feb 2020	Closed	Closed	Hive	S2 - Core Feature Inoperative	Murugesan Kumar		<p> 10 Feb 2020 Weekend data loading job was failed due to the containers are crashed. Checked the LLAP logs and identified Caused by: java.lang.IllegalArgumentException: No running LLAP daemons!</p> <p>Shared the llap logs for the last weekend.</p> <p> 19 Feb 2020 Mariano identified that yarn node manager crashed due to the llap memory: In memory cache per daemon to 144GB from 148GB.</p> <p>Post that we ran multiple runs and didnt face the llap container crash. Post confirmation close the case.</p>	Unable to find the root cause in the logs.
181	657592	Dev Cluster Ambari Errors	 04 Feb 2020	Closed	Resolved	Ambari (at least)	S2 - Core Feature Inoperative	Phil Young	Phil Young	<p>PROBLEM: While installing Druid Ambari was unable to restart services</p> <p>It was discovered that at 10:21 (while installing Druid) the Kerberos config parameters in all config files were removed.</p> <p>SOLUTION: We reverted the config for every service and restarted the services.</p>	Post deletion of incomplete Druid installation services are UP

180	657164	Setup Data Analytics Studio with Knox SSO	 03 Feb 2020	Closed	Resolved	knox	S2 - Core Feature Inoperative	Murugesan Kumar	Saurab	<p>KNOX is installed in our cluster. But we need support to install DAS in our HDP 3.1 cluster. Our cluster is connected with RedHat repository server which needs packages to be installed through repo files.</p> <p> 06 Feb 2020 Downloaded the DAS lite package and waiting for Support team to install in PERF cluster</p> <p> 07 Feb 2020 Cloudera support shared the linkg</p>	Waiting for the Support team to provide the solution
179	656108	Install Druid on Dev Cluster.	 03 Feb 2020	Closed	Resolved	Druid	Assistance from Cloudera to install in g component	Murugesan Kumar	Phil Young	<p>Assistance from Cloudera to install Druid.</p> <p> 06 Feb 2020 installed the Druid in masternode2/3 as per the guidelines from Cloudera support. Waiting for DEV team to validate the same.</p> <p> 07 Feb 2020 Team validates and reported GSS exception. Shared the comments with Cloudera support to get solution on this.</p>	Installed in DEV cluster. Unable to run the queries.



178	655499	Spark streaming jobs are getting failed after	 27 Feb 2020	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Gopinath Nagarajan	<p>java.io⁵.IOException: org.apache.hive.streaming.TransactionError: Aborted transaction cannot be committed: Transaction txnid:61291119 already aborted</p> <p>Spark jobs are getting failed and reported GSS exception in the logs.</p> <p>There was timeout exception in the application logs for Kerberos.</p> <p>Investigated that krb5.conf file getting over write in the cluster nodes. Corp IT team suggested to uncheck the Manage the krb5.conf file when we restart the Ambari services.</p> <p>Post that the spark jobs are running fine now.</p>	IT support fixed this issue along with the Cloudera support.
177	654172	Migrate the Hive data owner from perf-hive to perfhive	 21 Jan 2020	Closed	resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar		<p>This is the sub-case of #644677. The new user "perfhive" created in the LDAP server. Need Cloudera support on this to migrate the same.</p> <p>had WebEx with the support team and found that logs are getting permission issue and fixed to bring up the metastore, hiveserver2 and llap0.</p> <p>llap1 still fails to start on the container nodes. shared the logs to fix the issue</p> <p>llap1 has permission issues in /var/run directory</p>	Cloudera support team fixed and brought up LLAP1. case closed with comments



⁵ <http://java.io>

176	653597	Performance issues for Hive Queries	 17 Jan 2020	Closed	Closed	Spark	S2 - Core Feature Implementation	Murugesan Kumar	Pablo Jung	<p>getting vectorization issues when we run the hive queries, multiple count distinct with a group-by key.</p> <p>HWX confirmed as BUG for this query https://issues.apache.org/jira/browse/HIVE-22448</p> <p>Shared the export data and table creation script with them to analyse on that.</p> <p>Applied the DEV jar in PERF cluster and validation in progress.</p> <p> 05 Feb 2020 there are query performance issues reported. Cloudera confirmed to be checked by Engineerign team on their end.</p> <p> 07 Feb 2020 Developer reported exception for this patch: (Profile -Top 100 transactions by volume/Value) which is using DENSE_RANK() function in the queries.</p> <p>Updated the error with Cloudera</p>	Waiting for support team to release the patch for this BUG reported. DEV jar released.
-----	--------	-------------------------------------	---	--------	--------	-------	----------------------------------	-----------------	------------	---	--



175	64467 ⁶	Unable to run the Spark-submit getting permission denied exception for /tmp path	 03 Dec 2019	Closed	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Pankaj Kadam	<p>Unable to run the spark submit in PERF cluster. Works fine with the DEV cluster. Streaming jobs are fine. Only batch jobs are getting failed with the following exception: Caused by: org.apache.hive.service.cli.HiveSQLException: Error while compiling statement: FAILED: HiveAccessControlException Permission denied: user [srvredi] does not have [READ] privilege on [hdfs://PERFNNHA/tmp/20191203092354-515600dc-3453-4c03-82ab-6ed44b301962]</p> <p>Had multiple WebEx sessions and done apple to apple comparison. Identified that perf-hive user causes the issue.</p>	<p>HWX support proposed to create a new user called perfhive and validate.</p> <p>Waiting for downtime to make the changes.</p>
-----	--------------------	--	---	--------	-------	-------------------------------	-----------------	--------------	--	---

⁶ <https://my.cloudera.com/cases.html/644677>



174	651883	Unable to access the YARN Application logs which getting denied from yarn user too	 10 Jan 2020	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Murugesan Kumar	Akash Parmeshwar Sheldke	<p>Only Spark Streaming application logs unable to download. Getting exception:</p> <p>error: 20/01/10 01:30:47 INFO client.RequestHedgingRMFailover ProxyProvider: Created wrapped proxy for [rm1, rm2] 20/01/10 01:30:47 INFO client.AHSPProxy: Connecting to Application History server at</p> <p>There was the size of this application logs 175GB due that unable to download.</p> <p>Checked the config in spark code, custom log4j.properties with the spark-submit changed to WARN from INFO/Console.</p>	Post config changes the logs are reduced. Successfully completed the long run of the Streaming jobs.
173	651566	Hive LLAP is down and unable to bring up	 09 Jan 2020	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Magdum Sahib Sajjan	<p>There was performance tests we initiated last night. LLAP went down. Checked the LLAP logs and getting "no running LLAP daemons". Tried to bring up the LLAP services and couldn't bring up.</p> <p>Shared the llap logs and had WebEx with the support team.</p> <p>there was existing app hold the process. Killed through the perf-hive user in the masternode3 through:</p> <p>yarn app -destroy llap0 post that LLAP brought up</p>	Once destroyed the app, able to start the LLAP services.



172	650608	Spark Job fails with the error	 06 Jan 2020	Closed	Resolved	Spark	S2-Correlation Feature Inoperative	Murugesan Kumar	Suhashini Sukumar	<p>Spark submit fails with the error:</p> <p>Caused by: org.apache.hive.streaming.StreamingException: Unable to abort at org.apache.hive.streaming.HiveStreamingConnection\$TransactionBatch.close(HiveStreamingConnection.java:975)</p> <p>Shared the spark-submit code and the application logs too.</p> <p>The errors we encountered were due to a bug that caused excessive logging to std-out in the build of redi streaming. Now the jobs are running without any issues.</p>	Found error on spark code. Fixed by the dev team.
171	650564 ⁷	Hiveserver2 went down with same kerberos issues	 06 Jan 2020	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Kamireddy Kiran Reddy	<p>There is the same issue like #649787 for Hive services. Same Kerberos exception reported in Hiveserver2 logs.</p> <p>Had webex with Mariano, and found that krb5.conf file is there in both places. Updated the udp_preference_limit=1 in /etc/krb5.conf all the nodes of the cluster.</p>	As per the config provided, the services are looks good.



⁷ <https://my.cloudera.com/cases/650564/comments>


170	649787	Namenode service goes down 2 days once. Manual restart required	 30 Dec 2019	Closed	Resolved	HDFS	S2 - Core Feature Inoperative	Murugesan Kumar	Mohammad Naresh	<p>javax.security.auth.login.LoginException: Receive timed out Caused by: javax.security.auth.login.LoginException: Receive timed out</p> <p>Namenode service alone went down in the masternode1. ticket renew timeout exception reported in the namenode logs.</p> <p>Suggested to add, udp_preference_limit=1 in krb5.conf file in all the nodes of the cluster.</p>	As per the config provided, the services are looks good.
169	648225 ⁸	Error to retrieve the data on Hive table width 1000+ columns through Spark	 18 Dec 2019	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Nithin Nagaraj	<p>When we run the spark jobs for more than 1000+ columns getting this issue intermittently. There was workaround provided, - setting spark.sql.codegen.wholeStage=false. If the issue persists again will reopen new case</p>	NA

⁸ <https://my.cloudera.com/cases/648225/comments>

168	642110	Need to enable report metrics for Spark Streaming jobs.	 21 Nov 2019	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Anchal Chopra	PERF engineer wants to capture the reports through metrics. But we need to install the third party tool called Graphite and to configure the same through Grafana in Ambari.	Third party tool can't be possible and asked to close the case.
167	HWX# 00245115633332	Enabled Phoenix but unable to get the service in Hbase	 30 Oct 2019	Closed	Resolved	HBase	S2 - Core Feature Inoperative	Murugesan Kumar	Mayank Gupta	Have HDP 3.1 and installed Hbase as single Master and Region server for POC purposes. In the configuration page we enabled Phoenix services. but the summary page we don't have phoenix services to bring up. in hbase-site.xml added phoenix.schema.isNamespaceMappingEnabled=true phoenix.schema.mapSystemTablesToNamespace=true	Once added in the configuration, able to bring up the Phoenix services.

166	HWX#00244276	Unable to run 'CREATE TRIGGER' query in beeline	 23 Oct 2019	Closed	Resolved	Hive	S2-Core Feature Inoperative	Murugesan Kumar	Rohit Rami Mahottra	<p>Unable to run CREATE TRIGGER query in hive beeline and getting error. HWX confirmed when we done the upgrade from HDP 2.6.5 to 3.1, getting this metastore error. BUG-118520</p> <p>Metastore alter commands fixes this bug manually.</p>	Once done the alter command in Hive metastore, able to run the Create Trigger query in beeline.
165	HWX#00244274	LLAP status issue and Jar localization failing during LLAP startup	 23 Oct 2019	Closed	Resolved	Hive	S2-Core Feature Inoperative	Murugesan Kumar	Kunal Bhattacharya	<p>For POC Speedpay cluster unable to start Hive LLAP.</p> <p>The application status is checked through random ephemeral ports and this is defined as a range. These ports were blocked by the firewall</p> <p>Advised to set the ephemeral port range as below and make these ports available through the firewall: net.ipv4.ip_local_port_range = 32768 60999</p>	Once the required ports are opened able to start the Hive LLAP

164	HWX#0024273	YARN Timeline Service 2.0 and Resource Manager startup failed	 23 Oct 2019	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Kunal Bhattacharya	<p>For POC Speedpay cluster unable to start the timeline services due to memory inconsistent issues.</p> <p>We resolved the YARN Timeline Service 2.0 by checking <code>is_hbase_system_service_launch</code> and the service came up successfully.</p>	The services are UP and running.
163	HWX#0024087	Unable to view spark application logs in RM UI - HTTP ERROR 500	 22 Oct 2019	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Murugesan Kumar	Vikas Dadhich	<p>Spark UI application logs from the RM UI with the valid Kerberos ticket. can able to see rest of the application logs through that.</p> <p>Checked the cacerts file and looks fine.</p> <p>Remove the option for <code>spark.ssl.ui.enable</code> from Ambari it will be true by default if <code>spark.ssl.enable</code> is true</p>	Once disabled in Ambari - Spark conf, able to access the logs.

1 6 2	H W X# 00 24 27 69 6 3 3 3 2 3	Unable to run the 'Create Resource Plan' getting permission denied exception in DEV cluster	 10 Oct 2019	C o s e d	Res olv ed	Ra ng er	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Me en u Ku m ari	For Performance tuning purposes, need to create resource plan for Hive which gets failed with permission denied exception. Tried with svcredi/hive user with the appropriate privileges provided. Truststore/keystore for the ranger admin was incorrect. HWX support person fixed the same.	No w abl e to cap ture the res our ce pla n.
1 6 1	00 24 22 50 ⁹	Need help to install HUE in HDP 3.1 cluster	10/7/ 2019 2:14 AM	C l o s e d	Res olv ed	hiv e	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Gu lsh ad	Data Science team requires HUE tool. Checked in Ambari which is not available. Raised the case with HWX and confirmed that, HUE is not supported in HDP 3.1.x versions and suggested to use Ambari - files view which is a alternative solution	Con figu red the sam e for DS tea m and sha red the req uire d det ails.

⁹ <https://support.hortonworks.com/s/case/5002R000014ujl9QAA/detail>

160	00242224 ¹⁰	Configure the LLAP along with Node Lables	10/6/2019 10:30 PM	Closed	Resolved	hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Gunal	SpeedPay team needs new POC cluster deployment. Enabled the LLAP services which fails. Raised the case to configure the same with Node labels.	4 nodes POC cluster which needs to configure the llap services.
159	00242200 ¹¹	Unable to run the "Create Resource Plan" getting permission denied exception	10/6/2019 6:11 AM	Closed	Resolved	hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Shobika	<p>CREATE RESOURCE PLAN high_concurr_rp WITH QUERY_PARALLELISM=15; Error: Error while compiling statement: FAILED: HiveAccessControlException Permission denied: user [hive] does not have [SERVICEADMIN] privilege on [cov3lhdpms01.am.tsacorp.com¹²] (state=42000,code=40000)</p> <p>In Ranger we have given sufficient permissions for the 'srvredi' user. Shared the hiveserver2 logs and Ranger Admin logs too.</p> <p>there is ranger-admin-keyadmin.jks file hasn't been sync in all the masternodes. Done the same from ranger admin node. Now able to run the same command with 'srvredi' user.</p>	There is cluster ssl security has n't sync. Due that the ranger policy has n't been updated.

¹⁰ <https://support.hortonworks.com/s/case/5002R000014uisxQAA/detail>

¹¹ <https://support.hortonworks.com/s/case/5002R000014ufn3QAA/detail>

¹² <http://cov3lhdpms01.am.tsacorp.com/>

158	00241923 ¹³	JVM Threads Waiting continually increase on LLAP nodes over long-running tests	10/3/2019 4:40 AM	open	Pending with inputs to be shared	Hive	S2 - Core Feature Incorporated	Murugesan Kumar	Maria no	<p>Perf team reported with the IBI reporting tool have a potential problem waiting JVM threads continuously increasing in number across all nodes.</p> <p>HWX needs to capture the jstack inputs.</p> <p>Already shared the jstack alone but HWX needs Daemon logs too when we take the jstack</p>	Cluster is busy for continuous long run jobs. will capture the jstack and daemon logs
-----	------------------------	--	-------------------	------	----------------------------------	------	--------------------------------	-----------------	----------	--	---

¹³ <https://support.hortonworks.com/s/case/5002R000014uKVIQAM/detail>

157	00240389 ¹⁴	Unable to upgrade HDP 3.1.0.158-1	9/20/2019 8:33 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Maria no	<p>PERF Cluster repo issues reported while upgrading the HDP 158-1 hotfix.</p> <p>Installed the old packages due to the old repo configured. Removed the packages and configured the latest required repo in all the nodes of the cluster.</p> <p>Also deleted the metadata entries of the old packages.</p>	Removed the old repo in all the nodes for all the services too. now able to complete the HDP upgrade through express mode.
-----	------------------------	-----------------------------------	-------------------	--------	----------	--------	-------------------------------	-----------------	----------	--	--

¹⁴ <https://support.hortonworks.com/s/case/5002R000014Tn2qQAC/detail>

1 5 6	00 23 99 42 ¹⁵	Unable to access hive DB with Active-Passive LLAP HA JDBC URL	9/18/2019 12:13 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Maria no	Unable to connect hive DB with the HA JDBC url. HWX replied that this feature is not available in this HDP 3.1.x versions. Suggested to connect the Hive DB through dedicated LLAP url instead of zookeeper string	BUG-103199 This same problem was reported by ACI (Mithun Polisetty) via case 00229455 in early July .
-------------	------------------------------------	---	--------------------	--------	----------	------	-------------------------------	-----------------	----------	---	---

¹⁵ <https://support.hortonworks.com/s/case/5002R000014TZ43QAG/detail>

155	00239456 ¹⁶	Controlling Spark application log level in the driver and executor	9/13/2019 12:08 PM	Closed	Resolved	Spark	S2-Correlation Feature Inoperative	Murugesan Kumar	Akash	Controlling Spark application log level in the driver and executor. We are able to control spark Driver logs not the executor logs.	Team fixed to tune the spark-submit to control the logs.
154	00238616 ¹⁷	Exceptions reported in Hive Metastore logs :AlreadyExistsException	9/9/2019 12:36 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Karthik	There are multiple exceptions in the Hive metastore logs when we run the transformation and ingestion longrun jobs. there is similar case reported #00238097 to separate this exception to be addressed through separate case. Error: 2019-09-03T12:59:05,864 ERROR [pool-8-thread-186]: metastore.RetryingHMSHandler (RetryingHMSHandler.java:invokeInternal(201)) - AlreadyExistsException(message:Partition already exists:	HWX suggested that , this particular issue is harmless and can be safely ignored.

¹⁶ <https://support.hortonworks.com/s/case/5002R000014TE5GQAW/detail>

¹⁷ <https://support.hortonworks.com/s/case/5002R000014SjIBQA0/detail>

153	002380 ¹⁸	Getting FileNotFoundException in metastore logs	9/4/2019 6:44 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Tamil	<p>Errors reported from the metastore logs</p> <p>Error 1:</p> <p>FileNot Found is only for 6 tables:</p> <p>Solution:</p> <p>The above error will be printed when the table is not dropped from Hive but the HDFS directory for the table is removed from HDDS directory. Either you have to recreate all the partitioned directories and if the table is not required (as name says bkp) you can drop the table from hive.</p>	removed the tables from the DB which fixed this issue.
152	002376 ¹⁹	Unable to start LLAP after upgrade	8/30/2019 7:42 AM	Closed	Duplicate	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Juan Chietta	closing this case as duplicate case of 00237678	

¹⁸ <https://support.hortonworks.com/s/case/5002R000014SEo3QAG/detail>

¹⁹ <https://support.hortonworks.com/s/case/5002R000014RoZoQAC/detail>

151	00237678 ²⁰	Unable to start the Hive LLAP services in HDP upgrade	8/30/2019 7:39 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Hernandez Retamal	<p>When we tried to do Express upgrade facing issues to start the Hive LLAP services.</p> <p>LLAP was not starting, mostly related it cannot communicate with YARN resource manager. Yarn was failing because was unable to connect to Ranger KMS.</p> <p>Ranger ISSUE: Ranger Admin was started but the UI not working, this is related to the absence of the keystore file. We copied the keystore from a different location. The same occurred with KMS.</p> <p>After that Ranger > Yarn and LLAP was working fine</p> <p>Finally the LLAP services are brought up.</p>
150	00237652 ²¹	Unable to start Ranger Admin services to complete the HDP upgrade	8/30/2019 3:14 AM	Closed	Resolved	Ranger	S2 - Core Feature Inoperative	Murugesan Kumar	Tejas Parabhat	<p>When we tried to do Express upgrade facing issues to start the Ranger Admin services.</p> <p>package hasn't installed properly. So reinstalled the packages which caused the issues.</p> <p>Created conf folder under below and restarted both services</p> <p>/usr/hdp/current/ranger-usersync/</p> <p>/usr/hdp/current/ranger-tagsync</p>

²⁰ <https://support.hortonworks.com/s/case/5002R000014RoyZQAS/detail>

²¹ <https://support.hortonworks.com/s/case/5002R000014RnIXQAS/detail>

1 4 9	002374 ²²	Unable to bring up the LLAP in HDP patch upgrade	8/29/2019 4:18 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Hernandez Retamal	<p>We have received the hotfix for the HDP 3.1.0.139-1 for the other issues. Installed in all the nodes for this repo. When we try to do the express installation, couldnt able to start the LLAP services.</p> <p>On the first one there was a permission issue on /var/ambari-agent/tmp and not enough disk space on /tmp. After clean up the disk LLAP started ok On the second node, the symlink of /usr/hdp/current was not pointing to the right version, after fix that, both nodes restarted ok and the upgrade moved to the SPARK section.</p>	
1 4 8	002360 ²³	We are seeing HWC connector writing as NULL value after the New DEV JAR patch applied	8/19/2019 9:08 AM	Closed	Resolved by customer	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Pallo	<p>We applied the latest Hive warehouse connector jar in perf cluster and we get the following issue:</p> <p>We are seeing HWC connector writing as NULL value, even If we populate field value correctly, this behavior observed in 'Stream To Stream' option alone and when no of field is more than 400. Tested with upto 300 its working fine, as of now 400+ is breaking point.</p> <p>Support team still analyzing the issue.</p> <p>Issue has been identified, when schema has special char or value has special char its breaking, fixed now. please close the case</p>	


²² <https://support.hortonworks.com/s/case/5002R000014RhBGQA0/detail>

²³ <https://support.hortonworks.com/s/case/5002R000014QxEEQA0/detail>

1 4 7	00235819 ²⁴	Unable to complete the HDP Patch upgrade - HDP-3.1.0.126-1	8/16/2019 8:57 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Heenan	<p>In PERF cluster, we are not able to install the latest repo in only one node. It was still getting the old installation files in the masternode3:</p> <pre>[root@cov3lphdpmst03 ~]# yum repolist grep HDP-3.1.0.126-1</pre> <pre>!ACI_CorplT_HortonWorks-HDP_HDP-3_1_0_126-1 HDP-3.1.0.126-1 201</pre> <p>HWX support engineer fixed the repo files and completed the upgrade process.</p>	
1 4 6	00234910 ²⁵	Unable to install HDP-3.1.0.126 from the registered version page	8/9/2019 2:21 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Aashish	<p>Registered versions are not able to view in the Ambari's upgrade page. Problem identified in the postgres for DB.</p> <p>Fixed the same through the following instructions by HWX support engineer:</p> <pre>-- su - postgres</pre> <pre>-- psql -U ambari ambari</pre> <pre>Password : </etc/ambari-server/conf/password.dat> <ambari user database password> --</pre> <pre>ambari=> SELECT cluster_id, upgrade_id, cluster_name FROM clusters;</pre> <p>Post changes now upgraded the HDP 3.1.0.126-1 version in the cluster.</p>	

²⁴ <https://support.hortonworks.com/s/case/5002R000014Qq46QAC/detail>

²⁵ <https://support.hortonworks.com/s/case/5002R000012rglKQAQ/detail>

1 4 5	00234759 ²⁶	LLAP Interactive logs threw warnings which has to be addressed	8/8/2019 4:25 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Krishnadevan	<p>getting unwanted warning messages reported in the Hive Interactive logs. Even though those are warnings. But team needs to address the messages for the clarity on that.</p> <p>HWX support suggested to add: datanucleus.autoStartMechanism Mode=Ignored in hive-site.xml and hivemetastore.xml files. Post changes we are not getting this warning messages.</p>	
1 4 4	00234501 ²⁷	LLAP Containers are unstable	8/6/2019 9:47 PM	Open	Pending	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Paablo	<p>When we run the IBI concurrent users test, the llap containers are going down one by one. Reported to HWX as BUG. They released the DEV jar which works fine. Need to get the official patch for the same.</p> <p>Spark jobs performance issues reported.</p> <ol style="list-style-type: none"> 1.HWC run the df.count, query lands on hiveserver2Interactive, get the split 2. Driver distributed the splits to executors 3. Executors created the LLAP tasks based on split and submitted to llap 4. upon receiving these app tasks, LLAP failed the task with HDFS error 5. Llap client running in the spark executor is waiting forever for this data to be available <p> 30 Aug 2019 Engineering engaged to fix this case.</p>	

²⁶ <https://support.hortonworks.com/s/case/5002R000012rap2QAA/detail>

²⁷ <https://support.hortonworks.com/s/case/5002R000012rSZwQAM/detail>

1 4 3	00 23 19 26 ²⁸	Unable to bring up the Hive LLAP services	7/18/2019 2:52 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Javier	As part of HDP upgrade, the process gets failed due to LLAP services unable to start. Identified that node labels are causing the issue. HWX support engineer helped us to remove the placement_policy in the templates.py file which was holding the LLAP startup. Once removed the same we are able to start the LLAP services.	
1 4 2	00 23 19 19 ²⁹	Unable to start the Spark thrift server	7/18/2019 1:29 AM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Hari Kumar	In POC cluster unable to bring up the Spark thrift server due to resources issue in the default queue configured. Changed the queue in the spark configurations by HWX engineer. Now the Spark thrift server brought up to complete the upgrade process.	

²⁸ <https://support.hortonworks.com/s/case/5002R000012q5UmQAI/detail>

²⁹ <https://support.hortonworks.com/s/case/5002R000012q5J1QAI/detail>

141	00231342 ³⁰	Unable to bring up the hive LLAP post restart the Hive services.	7/15/2019 1:56 AM	Closed	Resolved by Customer	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nisha	There was config change made and tried to restart the Hive services which are getting failure. Identified that, previous process hasnt been killed properly and delete the data in the zookeeper quorum. Brought up the services on our own.	
140	00230552 ³¹	vertex “file not found” issue in transformation job	7/8/2019 9:12 AM	Open	Support Analysis in Progress	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Pablo Junge	When we run the compaction, transformation job failed with delta folder not found exception. Reported with HWX along with the logs to fix the same.	

³⁰ <https://support.hortonworks.com/s/case/5002R000012pnaHQAQ/detail>

³¹ <https://support.hortonworks.com/s/case/5002R00001421IdQAI/detail>

1 3 9	00229777 ³²	Unable to bring up LLAP once enabled HA	7/1/2019 5:57 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Karthik Nadar	<p>Need to enable the Active-Passive LLAP HA. Followed by the guide https://docs.hortonworks.com/HDPDocuments/HDP3/HDP-3.1.0/managing-hive/content/hive_setup_multiple_hsi.html</p> <p>But both LLAP nodes are not able to start.</p> <p>identified that memory parameters are incorrect to start the LLAP services. HWX engineer fixed.</p>
1 3 8	00229263 ³³	Duplicates coming in to Hive tables using structured streaming	6/26/2019 3:54 AM	Closed	Support Analysis in Progress	Spark	S2-Correlation Feature Inoperative	Murugesan Kumar	Nitin Agaraj	<p>We find that the consumer which is running using structured streaming have generated duplicate records on few days but producer is only sending it once as confirmed by looking at the producer logs. All these duplicates have come on different dates and different times. In the consumer there is no aggregation happening in the consumer that might generate any duplicates.</p> <p>The checkpoint dir is also handled by spark and saved in HDFS.</p>

³² <https://support.hortonworks.com/s/case/5002R0000141ccPQAQ/detail>

³³ <https://support.hortonworks.com/s/case/5002R0000141IN7QAM/detail>

1 3 7	00229069 ³⁴	Having problem in re-processing any time period data from Kafka using structured streaming	6/25/2019 3:38 AM	Closed	Support Analysis in Progress	Spark	Shashank Kumar	Murugesan Kumar	Aaditya Gururaj	<p>We are trying to reprocess some transactions from kafka again for a day before, or in generic for any time period due some data corruption in our table. Could you please help us out on how to re-process the data from past from kafka using spark-structured streaming for some time-period.</p> <p>There are chances where we have to reprocess the data in our production scenario</p>	
1 3 6	00228565 (see page 3)	data mismatch from Hive and LLAP namespace	6/20/2019 2:26 AM	Closed	Support Analysis in Progress	Hive	Shashank Kumar	Murugesan Kumar	Deepak	<p>query output mismatch between normal HS2 and interactive URL. Shared the required logs with the output of the query too.</p> <p>Can't able to reproduce the error in DEV cluster. Also HWX engineer suggested to remove 'over' when we run the query</p>	

³⁴ <https://support.hortonworks.com/s/case/5002R000013PRSIQAO/detail>

1 3 5	00228366 ³⁵	Unable to add new node labels in YARN queue manager	6/19/2019 2:04 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Shantanu Singh	<p>To avoid the performance mismatch created the one more node labels for spark jobs. But unable to start the YARN queue manger once added the spark_nodes labels.</p> <p>HWX fixed it through configure the new lables upto each child level to update. Once done able to start the services.</p> <p>Now spark_nodes have 5 worker nodes (wrk01-05)</p> <p>llap_nodes have 5 llap nodes (llap01-05)</p>
1 3 4	00227551 ³⁶	Unable to bring up Hive LLAP once we configured Node-Labels in Yarn queue	6/12/2019 12:04 PM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Nuruddin Ahmed	<p>The same issue we faced when we upgraded the cluster from 2.6.5 to 3.1. There we temporarily disabled the node labels and brought up the LLAP service. It caused the LLAP daemons pick the container nodes randomly and face the performance slowness issues. Cloudera SA suggested to enable back.</p> <p>There is a workaround provided by HWX: https://community.hortonworks.com/content/supportkb/242438/unable-to-start-hive-llap-services-after-enabling.html</p> <p>As per this we removed the placement_policy from the templates.py file in all the nodes of the cluster. Now able to see the node lables in RM UI and configured the same in the BIG queue too. Now LLAP runs on the LLAP nodes which are configured in the node labels.</p>

³⁵ <https://support.hortonworks.com/s/case/5002R000013P4RSQA0/detail>

³⁶ <https://support.hortonworks.com/s/case/5002R000013Og8PQAS/detail>

1 3 3	00 22 69 72(see page 3)	Hive LLAP HA patch doesn't work from the hotfix #00213633	6/7/2 019 8:52 AM	C l o s e d	Sup por t Ana lysi s in Pro gre ss	Hiv e	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Bri an	LLAP HA patch raised by UAT team and applied the same in the POC cluster to validate the same. When we connected the LLAP url and able to run the queries and meantime when we kill the Active's LLAP daemon process ID the query got disturbed. After sometime we are unable to connect to the same JDBC url. Once we brought up the LLAP services only can resume the process through beeline. Shared the required logs to fix the same.	
1 3 2	00 22 69 19 ³⁷	LLAP Cache not being hit during the query runs even on multiple iterations.	6/7/2 019 4:04 AM	C l o s e d	Pen din g Eng ine erin g Inp ut	Hiv e	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Nis ha	30 CU tests unable to hold the cache and HWX engineer observed that LLAP cache status always be 0. Due that performance queries are taking more time to complete it. HWX will clone the new case for the Engg team to fix it.	

37 <https://support.hortonworks.com/s/case/5002R000013OMTRQA4/detail>

131	00226730 ³⁸	Even though LLAP is running, LLAP application is not showing in RM UI.	6/6/2019 2:27 AM	Closed	Resolved by Customer	YARN	S2-Correlation Feature Inoperative	Murugesan Kumar	Shobika	Unable to view the BIG queue's Jobs in the RM UI even though we have valid MIT kerberos keytab file. Checked with the performance team and able to access through 'srvredi' user account. Might be the issues with internal IT team. So asked HWX to close this case and start to access the RM UI with srvredi account.	
130	00226700 ³⁹	Unable to bring up LLAP services once yarn queues are re-configured - DEV cluster	6/5/2019 9:51 PM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Shobika	Duplicate case of #00226696 for DEV cluster. The same queue modifications are affected to bring up the LLAP daemons.	

³⁸ <https://support.hortonworks.com/s/case/5002R000013OHheQAG/detail>

³⁹ <https://support.hortonworks.com/s/case/5002R000013OH1kQAG/detail>

1 2 9	00226696 ⁴⁰	Unable to bring up LLAP services once yarn queues are re-configured	6/5/2019 9:51 PM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Shobika	Re-aligned the new queues are not sufficient for BIG queue to launch the daemon services. Reduced the no of LLAP containers are from 5 to 4 and increased the BIG queue size from 40% to 45%. Now able to start the services of LLAP daemons with 4 nodes of containers.	
1 2 8	00225444 ⁴¹	Unable to bring up Hive LLAP post Ambari upgrade 2.7.3.6-4	5/28/2019 3:13 AM	Closed	Resolved	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Karthik Nadar	Post upgrade the Ambari services unable to bring up the LLAP services alone in DEV cluster. HWX engineer identified that, the LLAP application hangs in the ACCEPTED mode. In YARN summary page, the nodemanagers are not running, due that the application waits for the resources to be allocated. Once we brought the nodemanager services able to bring up the LLAP services too.	

⁴⁰ <https://support.hortonworks.com/s/case/5002R000013OGqzQAG/detail>

⁴¹ <https://support.hortonworks.com/s/case/5002R000013NhfxQAC/detail>

1 2 7	00 22 54 32 ⁴²	Post upgrade unable to start the one of the Namenode and other services	5/28/2019 2:20 AM	Closed	Resolved by customer	HD FS	Spark Feature Inoperative	Mrugesh Kumar	Deepak Chandr	As part of Ambari upgrade unable to start the Namenode. Identified in logs that existing process id holding the failure state of the Namenode services. Killed the process manually and brought the services up.	
1 2 6	00 22 44 59 ⁴³	Vertex error on Base table - delta folder not found	5/20/2019 9:04 AM	Closed	Analysis in progress	Hive	Spark Feature Inoperative	Mrugesh Kumar	Pablo Jung	<p>In Spark jobs we get delta folder not found exception and the jobs too get failed. Identified that the table Compaction doesnt run properly. The delta directories are there still and ready for cleaning state for a long while. Shared the logs to debug the same.</p> <p>HWX engineering team debugging the issue why the locks are there in the table.</p> <p>Received the official patch and validation is in progress now.</p>	

⁴² <https://support.hortonworks.com/s/case/5002R000013NhYDQA0/detail>

⁴³ <https://support.hortonworks.com/s/case/5002R000013NFoYQAW/detail>

1 2 5	00224033 ⁴⁴	Locking state always becomes WAITING when trying to truncate the Hlve table	5/16/2019 5:42 AM	Closed	Resolved	Hive	Se-Cor-Feat-ur-e In-operative	Murugesan Kumar	Karthik Nadar	<p>You will find that the other than truncate all the queries working and show locks; also shows that there is no lock pending on the table and after running truncate the Query gets hung and we need to abort the transactions to stop it. During this time if you see the locks you will find that the table was not able to acquire any lock and remains in waiting state.</p> <p>Won't be able to get the logs for you as It is not arising on the tables but if it occurs again (state lock situation where the Query has been executed and there are no current transaction on that table but still the truncate does not work or some other query does not work) then we shall collect all the logs.</p> <p>So the case is closed.</p>	
1 2 4	00223393 ⁴⁵	UPSERT operation not supporting to load the data into Hive.	5/12/2019 11:22 PM	Closed	Resolved	Spark	Se-Cor-Feat-ur-e In-operative	Murugesan Kumar	Anchal Chopra	<p>If Spark DataFrame contains both Update and new records are present then not able to load the data into Hive tables. for work around we are using Hive merge. Can you please let us know this feature is available in HDP 3.1 or not. If not can you please provide the patch.</p> <p>HWX confirmed that this feature is not available for this version:</p> <p>There is an open apache jira referring the issue: =====</p> <p>https://issues.apache.org/jira/browse/SPARK-19335</p>	

⁴⁴ <https://support.hortonworks.com/s/case/5002R000013N3mDQAS/detail>

⁴⁵ <https://support.hortonworks.com/s/case/5002R000013MoDKQA0/detail>

1 2 3	00223077 (see page 3)	Hive response slow for concurrent queries	5/9/2019 6:42 AM	Closed	Analysis in progress	hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nisha Vaidya	Hive performance are too slow when we run 30 concurrent users test. Uploaded the required logs to debug with HWX engineer.	
1 2 2	00222745 ⁴⁶	Unable to bring up the hive LLAP post restart the services.	5/7/2019 8:16 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nitin Garg	After restart the Hive services, unable to bring up the LLAP services alone in the PERF cluster. Identified that, the zookeeper quorum unable to clear the data to bring up the services. Deleted the data in quorum manually and brought the services through HWX support engineer. Ref: https://community.hortonworks.com/articles/29900/zookeeper-using-superdigest-to-gain-full-access-to.html https://community.hortonworks.com/content/supportkb/49531/authentication-failure-when-attempting-to-connect.html	

⁴⁶ <https://support.hortonworks.com/s/case/5002R000013MUNHQA4/detail>

1 2 1	00 22 21 31 ⁴⁷	Join results in ArrayIndex OutOfBoun dsExceptio n, when used executeQue ry() Patch not working HDP 3.1.0.22-3	5/1/2 019 9:47 PM	C l o s e d	Res olv ed	Sp ark	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Ra nji th Pu llu ru	There is a specific instruction from the Engineering team to use this patch: "spark.datasource.hive.warehouse.disable.pruning.and.pushdowns=true this property is hotfix and production ready too by HWX team.
1 2 0	00 22 08 47 ⁴⁸	Spark jobs are failed and nm/rm logs reported as yarn-ats permission denied issue	4/23/ 2019 4:28 AM	C l o s e d	Res olv ed	YA RN	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	De ep ak Ch an de r	Permission denied exception for yarn-ats user in the nn and rm logs. Cleared the data directory of the embedded hbase. After restarted the services, the issues got resolved.

⁴⁷ <https://support.hortonworks.com/s/case/5002R000013M8aqQAC/detail>

⁴⁸ <https://support.hortonworks.com/s/case/5002R000012RwAcQAK/detail>

119	00220708 ⁴⁹	Ambari not regenerating keytabs	4/22/2019 3:42 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Akash Shinde	<p>POC Cluster unable to proceed the Patch upgrade and reported error as unable generate keytabs for the services installed in the cluster.</p> <p>HWX identified that, the password used for admin principle was not correct so after changing the password to correct combination it completed the operation.</p>
118	00220684 ⁵⁰	Hiveserver2 goes down immediately	4/21/2019 10:47 PM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Balaji	<p>Hive goes down immediately even restart done manually through. Attached the HS2 logs to debug with HWX.</p> <p>Identified that, the previous running Hive process didn't stop yet. Killed the process manually and restarted the services. Now the Hive services are working stable.</p>

⁴⁹ <https://support.hortonworks.com/s/case/5002R000012RrZHQA0/detail>

⁵⁰ <https://support.hortonworks.com/s/case/5002R000012RrLeQAK/detail>

117	00220247 ⁵¹	Installation failed for Zeppelin service	4/16/2019 11:38 PM	Closed	Resolved	Zeppelin	S2 - Core Feature Inoperative	Murugesan Kumar	Akash Shinde	<p>In POC Cluser - zeppelin installation failed for Data Science team.</p> <p>Identified the service user configured as 'zeppelin' instead of poc_zeppelin. HWX re-configured and able to bring up the services now.</p>	
116	00220243 ⁵²	Structured streaming Job failed after some hours of run with no particular ERROR seen in the logs	4/16/2019 11:10 PM	Closed	No Customer Response	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Vikas Daidich	<p>In the logs we repeated saw two kinds of warnings: 19/04/16 13:16:50 INFO SchemaManager\$: Trying to get schema for subject 'null' and id '21' 19/04/16 12:01:28 INFO Fetcher: [Consumer clientId=consumer-2, groupId=spark-kafka-source-8a4465cb-483e-4035-b7c1-36f42e01aea6--1454838663-driver-0] Resetting offset for partition perf-exec-7-0 to offset 152200167.</p> <p>it is set 4GB while submitting the JOB driver is running out of the Heap.</p>	

⁵¹ <https://support.hortonworks.com/s/case/5002R000012RelwQAC/detail>

⁵² <https://support.hortonworks.com/s/case/5002R000012ReiTQAS/detail>

115	00219662 ⁵³	Vertex issue reported once we downgraded the Postgres	4/12/2019 12:11 AM	Client	Support Analysis in Progress	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nitin Garg	<p>Last Friday we were having, Postgres was psql (9.6.11, server 9.4.17). Due to some network issues, we restarted the nodes. Post restart, the Postgres data got crashed. Raised separate case #00218841 and Aashish brought down the postgres to psql (9.2.23) which we had dump file of last month got restored. Post fix, we are not able to run the same queries and gives vertex issues.</p> <p>Uploaded the required logs. HWX debugging in progress.</p>	
114	00218454 ⁵⁴	Spark Ingestion got failed after 18hrs continuous run	4/4/2019 4:04 AM	Client	Support Analysis in Progress	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Nidhin Divakaran	<p>We executed the Spark Ingestion job continuously for 18 hours and reported the error: Caused by: org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in stage 10347.0 failed 4 times, most recent failure: Lost task 0.3 in stage 10347.0 (TID 717402, cov3lphdpllap02.am.tsacorp.com⁵⁵, executor 1): java.io⁵⁶.IOException: Received reader event error: Timed out waiting for heartbeat for task ID attempt_7202709020297211622_89461_0_00_000000_3 Caused by: java.nio.channels.ClosedByInterruptException</p> <p>Uploaded the logs with HWX to debug the same. Analysis is in progress.</p>	

53 <https://support.hortonworks.com/s/case/5002R000012RPS2QAO/detail>

54 <https://support.hortonworks.com/s/case/50044000012QvQaAAK/detail>

55 <http://cov3lphdpllap02.am.tsacorp.com/>

56 <http://java.io>

1 1 3	00218198 ⁵⁷	Operation category READ is not supported in state standby	4/03/2019	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Murugesan Kumar	Ramakis hn a Reddy Kanchi	Getting WARN messages: WARN [ContainerLocalizer Downloader] org.apache.hadoop.ipc.Client: Exception encountered while connecting to the server : org.apache.hadoop.ipc.RemoteException(org.apache.hadoop.ipc.StandbyException): Operation category READ is not supported in state standby. Visit https://s.apache.org/sbnn-error Is it safe to ignore and how can we get rid of it? HWX suggested to do ignore those warnings	
1 1 2	217859 ⁵⁸	Official functionality available for doing the operation conversion from Schema type to sql Struct type	04-01-2019 06:25	Closed	Support Analysis in Progress	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Nithin Nagaraj	waiting for HWX confirmation to use the third party jar.	

⁵⁷ <https://support.hortonworks.com/s/case/50044000012Qnu5AAC/detail>

⁵⁸ <https://support.hortonworks.com/s/case/50044000012QfJvAAK/detail>

111	217608 ⁵⁹	Bytearray deserialization error while submitting spark-submit	3/29/2019 1:11 AM	Closed	Pending Customer Input	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Hari Kumar	Fixing is in progress. Requested SME support to fix this issue.	
-----	--------------------------------------	---	-------------------	--------	------------------------	-------	-------------------------------	-----------------	------------	---	--

⁵⁹ <https://support.hortonworks.com/s/case/50044000012QWLVA4/detail>

110	217593 ⁶⁰	Unable to do patch upgrade - Timeline Service V2.0 Reader couldnt start	3/28/2019 9:43 PM	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Murugesan Kumar	Aashish Bhatla	<p>1. We have found the bellow error in logs;</p> <p>"java.io⁶¹.FileNotFoundException: file:/usr/hdp/3.1.0.0-78/hadoop-yarn/timelineservice/hadoop-yarn-server-timelineservice-hbase-coprocessor.jar⁶²"</p> <p>2. We removed the ats-hbase data from below command. # hdfs dfs -ls /atsv2/hbase/data</p> <p>3. We also removed the Remove the zookeeper znode /atsv2-hbase-secure as it contains entries with acl set to the old user. # zookeeper-client >> rmr /atsv2-hbase-secure</p> <p>4. We also checked that multiple ats-hbase instance is running, So I removed the all instance and start the fresh ATSV2 Service from Ambari.</p>	This is configuration issue in our node(disk space issue). Deleted the old HDP 3.1 directory and caused the issue.
-----	--------------------------------------	---	-------------------	--------	----------	------	-------------------------------	-----------------	----------------	--	--

⁶⁰ <https://support.hortonworks.com/s/case/50044000012QVuJAAW/detail>

⁶¹ <http://java.io>

⁶² <http://file/usr/hdp/3.1.0.0-78/hadoop-yarn/timelineservice/hadoop-yarn-server-timelineservice-hbase-coprocessor.jar>

109	217420 ⁶³	Exception in thread "spark-listener-group-eventLog" java.lang.OutOfMemoryError: Java heap space	3/27/2019 10:17 PM	Closed	Pending Customer Input	Spark	Spark - Core Feature Inoperative	Murugesan Kumar	Nidhin Divakaran(s ee page 3)	<p>We found that the Databricks Avro jar is not compatible with the spark version in HDP cluster and because of the jar Spark driver is consuming more heap space and finally exiting with OOM.</p> <p>In PERF cluster still fixing is going on.</p>	
108	217242 ⁶⁴	Unable to copy files to the Ranger KMS zone.	3/26/2019 10:37 PM	Closed	Resolved	Ranger KMS	Spark - Core Feature Inoperative	Murugesan Kumar	Kamireddy Kiran Reddy	<p>hadoop.kms.blacklist.DECRYPT_EK property - perf_hdfs user was added in Ambari.</p> <p>So it caused the authorization issue.</p> <p>Removed the user and restarted the Ranger KMS now able to copy the files into the KMS zone.</p>	Config issue. In Ambari for security issues.

⁶³ <https://support.hortonworks.com/s/case/50044000012QRXUAA4/detail>

⁶⁴ <https://support.hortonworks.com/s/case/50044000012QN1KAAW/detail>

107	217081 ⁶⁵	Zookeeper servers are inconsistent	3/26/2019 4:27 AM	Closed	Resolved	Zookeeper	S2-Correlation Feature Inoperative	Murugesan Kumar	Tauseef Hussain	<p>1) LLAP was failing to start due to below error: java.io⁶⁶.IOException: Packet len4194345 is out of range!</p> <p>2) We updated Zookeeper Client JVM Flags property and added jute maxbuffer size to be 100 MB.</p> <p>3) We stopped Hive, cleared out below znodes, and restarted Hive: /zkdtsm_hive_llap0/ZKDTSMRoot/ZKDTSMTokensRoot /hive/cluster/delegationMETASTORE/tokens</p> <p>4) After this LLAP starts up fine and we validated that queries run fine.</p> <p>Also both cluster have interconnected through the DEV's hbase-site.xml in PERF cluster. Moved from spark conf directory and Zookeeper nodes are running fine as expected.</p>	Cluster config issues. By mistake copied the DEV's conf file in PERF cluster when they built.
106	216833 ⁶⁷	LLAP daemon is not stable restarted multiple times	3/24/2019 4:38 AM	Closed	Pending Customer Input	Hive	S2-Correlation Feature Inoperative	Murugesan Kumar	Rohit Rai	LLAP daemons are going down when we execute jobs. Fixing is in progress to stable the LLAP containers.	Observation is in progress

⁶⁵ <https://support.hortonworks.com/s/case/50044000012QI7tAAG/detail>

⁶⁶ <http://java.io>

⁶⁷ <https://support.hortonworks.com/s/case/50044000012QAOWAAO/detail>

105	216491 ⁶⁸	Spark jobs failing with Memory leak issue in Hive	3/20/2019 11:37 PM	Closed	Duplicate	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Vikas Dadhich	This is same as "ArrayOutOfBound Exception" and HWX provided the patch HDP 3.1.0.22-. Validation is in progress in POC cluster.	Duplicate
104	216302 ⁶⁹	Unable to launch Knox gateway	3/20/2019 12:45 AM	Closed	Resolved	Knox	S2 - Core Feature Inoperative	Murugesan Kumar	Kamireddy Kiran Reddy	Knox gateway Admin UI unable to start due the incorrect credentials in Keystore. Updated the same now able to access the Knox UI (admin/admin-password). But the IBI UI couldnt configure through Knox. It supports only Hadoop services alone.	IT team misplaced the credentials and recreated.

⁶⁸ <https://support.hortonworks.com/s/case/50044000012PzGIAAK/detail>

⁶⁹ <https://support.hortonworks.com/s/case/50044000012Pt2GAAS/detail>

103	216282 ⁷⁰	Unstable environment, datanodes are going down in few minutes post HDP upgrade	3/19/2019 9:52 PM	Closed	Resolved	HDFS	S2 - Core Feature Inoperative	Murugesan Kumar	Ramakrishna Reddy Kanchi	Due to millions of files in blocks. The default Heap for Datanodes are 1GB and increased to 4GB and now the datanodes are stable in PERF cluster.	Configuration issue when they built the cluster. The data growth has come now.
102	216125 ⁷¹	Unable to do patch upgrade from HDP 3.1.0.0-78 to HDP 3.1.0.10-1	3/19/2019 3:35 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Pankaj Kadam	Chose Rolling upgrade option and it hangs in "Pre Upgrade Hive Backup" step. Even we had taken the metastore backup manually in postgres node. Unable to proceed for the patch upgrade from HDP 3.1.0.0-78 to HDP 3.1.0.10-1. cache was holding the ambari server. Paused the upgrade and restarted the Ambari server. Then the we got the step to select the back up stage. Successfully installed the patch in the PERF cluster.	Ambari upgrade issue caused the issue.

⁷⁰ <https://support.hortonworks.com/s/case/50044000012PsdUAAS/detail>

⁷¹ <https://support.hortonworks.com/s/case/50044000012Pnz9AAC/detail>

101	215932 ⁷²	Spark shell - Hive Warehouse connector query different output for executeQuery	3/18/2019 5:45 AM	Closed	Resolved	Spark	S2 - Core Feature Incorporated	Murugesan Kumar	Ranjith	<p>BAE team reported:</p> <p>In spark shell while using executeQuery val df = hivesession.executeQuery("select OID from redi.trans_master_core_utc where OIDateYYYYMMDD >= 20180301 and OIDateYYYYMMDD <= 20190310").show(100, false) Result: 1 record</p> <p>while using execute val df = hivesession.execute("select OID from redi.trans_master_core_utc where OIDateYYYYMMDD >= 20180301 and OIDateYYYYMMDD <= 20190310").show(100, false) Result : 23 records</p> <p>But HWX and BAE both are not able to reproduce this error and asked HWX to close this case.</p>	Not able to reproduce this issue.
100	215930 ⁷³	Hive compaction not working.	3/18/2019 5:31 AM	Closed	Pending Customer Input	Hive	S2 - Core Feature Incorporated	Murugesan Kumar	Nitin garg	<p>Fixing is still going on.</p> <p>Can you run the compaction like below:</p> <pre># set mapreduce.map.memory.mb=16384; # set mapreduce.map.java.opts=-Xmx12288m; # alter table redi.rule_hits_ingestion partition(oibaseyyyyymmdd=20190311) compact 'minor';</pre>	Still fixing is in progress. HDP 3.1 issue

⁷² <https://support.hortonworks.com/s/case/50044000012PjXIAAK/detail>

⁷³ <https://support.hortonworks.com/s/case/50044000012PjVAA0/detail>

99	215675 ⁷⁴	Incorrect Timestamp from Spark-executeQuery	3/14/2019 11:34 PM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Pankaj	<p>facing issues with executeQuery(), where in, looks like its converting the timestamp which is in UTC in hive table to system time zone timestamp(CST). But execute() and from beeline it returns the timestamp as it is in hive table.</p> <p>HWX gave solution as to add this property in the spark-shell and check whether you are seeing same values: ++++++ --conf "spark.driver.extraJavaOptions=-Duser.timezone=UTC" --conf "spark.executor.extraJavaOptions=-Duser.timezone=UTC"</p>	HDP 3.1 issue
98	215278 ⁷⁵	Unable to access Superset with valid Kerberos credentials	3/13/2019 1:56 AM	Closed	Resolved	Superset	S2 - Core Feature Inoperative	Murugesan Kumar	Harikumar D	<p>Unable to find the Hive view in HDP 3.1 and Superset has introduced. Not able to start the same.</p> <p>HWX suggested that Superset is still in progress and need to wait for the next stable version.</p>	HDP 3.1 issue

⁷⁴ <https://support.hortonworks.com/s/case/50044000012PbMnAAK/detail>

⁷⁵ <https://support.hortonworks.com/s/case/50044000012PRF1AAO/detail>

97	214473 ⁷⁶	Patch upgrade for HDP 3.1	03-06-2019 21:16	Closed	Resolved	Ambari	S2-CORE Feature Inoperative	Murugesan Kumar	Deepak Chander	Observed that /usr/hdp/current directory still has the old jars and cause this issue. HWX suggested to do the soft link for the old jars. Done manually and working as expected.	HDP 3.1 issue
96	214459 ⁷⁷	Unable to write dataframe into Hive from spark	03-06-2019 18:45	Closed	Resolved	Spark	S2-CORE Feature Inoperative	Murugesan Kumar	Anchal Chopra	hive.load.data.owner was there as hive and we have user called perf-hive in PERF cluster. Changed the data owner to perf-hive from hive and working as expected.	HDP 3.1 issue

⁷⁶ <https://support.hortonworks.com/s/case/50044000012P5YcAAK/detail>

⁷⁷ <https://support.hortonworks.com/s/case/50044000012P5D8AAK/detail>

95	214295 ⁷⁸	Unable to write dataframe into Hive from spark	03-06-2019 01:46	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Rishabh Shukla	Duplicate of the case to write data into Hive from Spark through Hive Warehouse Connector	Duplicate
94	213928 ⁷⁹	Non Availability jars in HDP3.1	03-04-2019 06:56	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Jagadeesan AS	Below jars are missed in HDP 3.1 but available in HDP 2.6.5 1.kafka-clients-1.0.0.2.6.5.0-292.jar 2.spark-streaming-kafka-0-10_2.11-2.3.0.2.6.5.0-292.jar HWX suggested to download from their repository to use this JAR.	suggestion asked from HWX

⁷⁸ <https://support.hortonworks.com/s/case/50044000012P1hhAAC/detail>

⁷⁹ <https://support.hortonworks.com/s/case/50044000012OsmNAAS/detail>

93	213668 ⁸⁰	Hive Vertex issue for queries	03-01-2019 03:25	Closed	Support Analysis in Progress	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Kartik Nadar	<p>HWX accepted that, there is a BUG and released the Dev JAR to fix this issue.</p> <p>for workaround:</p> <p>set hive.vectorized.execution.enabled=false;</p> <p>fixing in progress</p>	HD P 3.1 issue
92	213534 ⁸¹	Unable to bring up Ranger Admin and Ranger client services	2/28/2019 9:39 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Ricardo Fuentes	<p>As part of PERF cluster network outages, reboot all the nodes. And couldnt able to bring up the services.</p> <p>After checking two nodes, the postgresql database server was down. We get it up, and every service on the cluster is working now.</p>	Config issue

⁸⁰ <https://support.hortonworks.com/s/case/50044000012OjM3AAK/detail>

⁸¹ <https://support.hortonworks.com/s/case/50044000012OfwCAAS/detail>

91	213244 ⁸²	Select Hive query running slow.	2/27/2019 1:46 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nitin Garg	Nitin requested to open separate case for select query from #213241	Duplicate
90	213241 ⁸³	Hive Query takes very slow for update/select	2/27/2019 1:38 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nitin Garg	hive.compactor.max.num.delta=100 hive.tez.container.size=8GB	HD P 3.1 issue and config issue

⁸² <https://support.hortonworks.com/s/case/500440000120aHJAA0/detail>

⁸³ <https://support.hortonworks.com/s/case/500440000120aGkAAK/detail>

89	212808 ⁸⁴	Update Hive query causes Java Heap space issue	2/25/2019 1:39 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar	Nitin Garg	were able to run the query after changing the hive.tez.container.size to 8 GB.	config issue
88	212771 ⁸⁵	Join results in ArrayIndexOutOfBoundsException, when used executeQuery()	2/24/2019 9:16 PM	Closed	Fixing in progress	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Pankaj	Identified as BUG and they released the patch to fix this issue: HDP 3.1.0.22-3 ⁸⁶ patch installed in the POC cluster.	HDP 3.1 issue

⁸⁴ <https://support.hortonworks.com/s/case/50044000010u6ZZAAY/detail>

⁸⁵ <https://support.hortonworks.com/s/case/50044000010u6F0AAI/detail>

⁸⁶ <http://private-repo-1.hortonworks.com/HDP-GPL/centos7/3.x/updates/3.1.0.22-3/hdp.gpl.repo>

87	21 18 10⁸⁷	Add a new HDP Support portal User	2/18/2019 4:22 AM	Closed	Resolved	Admin	S2 - Core Feature Inoperative	Murugesan Kumar	Raksha Jayaraj	raised admin case to create new user account for Snehal and created too.	
----	---------------------------------------	-----------------------------------	-------------------	--------	----------	-------	-------------------------------	-----------------	----------------	--	--

⁸⁷ <https://support.hortonworks.com/s/case/50044000010tYvcAAE/detail>

86	210919 ⁸⁸	Post upgrade unable to start HiveInteractive services	02-11-2019 22:26	Closed	Resolved	Hive LLAP	S2 - Core Feature Inoperative	Murugesan Kumar	Kunal Bhattacharya	<p>ISSUE#1: -</p> <p>Hive LLAP was not coming up after upgrading from HDP-2.6.5 to HDP-3.1. Looking into the RM UI app logs we saw that it was failing due to the following error: Caused by:</p> <pre>org.apache.hadoop.yarn.exceptions.YarnException: Download and unpack failed at org.apache.hadoop.yarn.util.FSDownload.downloadAndUnpack(FSDownload.java:306) ... 10 more</pre> <p>Caused by:</p> <pre>org.apache.hadoop.hdfs.BlockMissingException: Could not obtain block: BP-18885567-10.5.26.236-1534212482606:blk_1081796646_21150566 file=/hdp/apps/3.1.0.0-78/yarn/service-dep.tar.gz - We checked for corrupt blocks at /hdp/apps/3.1.0.0-78/yarn/ and found that the file was corrupted - We ran fsck -delete command to remove the corrupted block - Copied the file service-dep.tar.gz from LocalFS to HDFS: \$ hdfs dfs -put /usr/hdp/3.1.0.0-78/hadoop-yarn/lib/service-dep.tar.gz /hdp/apps/3.1.0.0-78/yarn/ - After this we tried restarting Hive Interactive but it failed with unknown status error</pre> <p>ISSUE#2: - We found that there was Node Labels enabled on the cluster for LLAP - This is a known issue with HDP-3.1 where LLAP fails to start if Node Labels were enabled. - We disabled Node Labels from YARN Queue Manager and YARN configs and restarted LLAP - LLAP Service started without any issues and we also tested by running queries through Beeline.</p>	Config issue
----	--------------------------------------	---	------------------	--------	----------	-----------	-------------------------------	-----------------	--------------------	---	--------------

⁸⁸ <https://support.hortonworks.com/s/case/50044000010t6H6AAI/detail>

85	21 09 15 ⁸⁹	Post upgrade unable to start HDFS / YARN services through Ambari	02-11-2019 22:08	Closed	Resolved	HDFS	S2 - Core Feature Inoperative	Murugesan Kumar	Saurabh Lamba	<p>NN writes ranger-audits to infra-solr. When solr is not running, NN will write audits to local spool folder. The NameNode becomes very busy while writing audits locally and as a result it does not respond to the JournalNodes.</p> <p>The root cause for today's Namenode failure lies with NN not being able to communicate to JorunalNodes. We restarted all the 3 JNs and then tried starting NN back. This worked, and now both NN are online.</p>	Infrastructure issue
84	00 21 05 58 ⁹⁰	RESTART YARN/TIMELINE_READER failing	2/8/2019 8:22 AM	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Murugesan Kumar	Gulshad Ansari	RESTART YARN/TIMELINE_READER not starting up in the middle of upgrade from 2.6 to 3.1	Config issue

⁸⁹ <https://support.hortonworks.com/s/case/50044000010t6FAAAY/detail>

⁹⁰ <https://support.hortonworks.com/s/case/50044000010sw3hAAA/detail>

83	00210515 ⁹¹	Unable to RESTART RANGER while upgrade to HDP 3.1	2/8/2019 4:07 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Guishad Anshari	<p>We are in progress of HDP upgrade from 2.6.5 to 3.1. As part of upgrade activity, we couldnt restart the Ranger admin and services. Attached the ambari logs with ranger - tagsync logs herewith.</p> <p>HWX fixed this issue with keystore path:</p> <p>Created new dummy keystore # keytool -genkey -keyalg RSA -alias ranger-tagsync -keystore /etc/security/serverKeys/ranger-tagsync-keystore.jks -validity 360 -keysize 2048 -storepass hadoop #chown ranger:ranger /etc/security/serverKeys/ranger-tagsync-keystore.jks</p>	upgrade issues when HDP 2.6.5 to 3.1 in PERF cluster
82	00210311 ⁹²	Unable to view HDP 3.1 version in Ambari for Manager versions menu	2/7/2019 12:57 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Shantanu Singh	<p>We tried to register as new version for 3.1. But post install it shows, Stack HDP-3.1 Name HDP-3.1.0.0-78 Version 2.6.1.0-129 Can not able to delete and register as new version getting,</p> <p>HWX fixed this issue at Ambari Database and proceed for HDP upgrade</p>	upgrade issues when HDP 2.6.5 to 3.1 in PERF cluster

⁹¹ <https://support.hortonworks.com/s/case/50044000010svH8AAI/detail>

⁹² <https://support.hortonworks.com/s/case/50044000010sqrLAAQ/detail>

81	00210297 ⁹³	Spark failing to append data if table have existing data	2/6/2019 10:38 PM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Hariri	Already we raised the case for NULL insert #00209775. This case for Spark failing to append data if table have existing data. inputDataFrame.write.format(HIVE_WAREHOUSE_CONNECTOR).option("database", "redi").mode("Append").option("table", tableName).save() HWX fixed the issue.	This is related to 00209775 case.
80	00210105 ⁹⁴	Upgrade from HDP 2.6.5 to 3.1 Hive pre-upgrade error	2/6/2019 3:08 AM	Closed	Resolved	hive	S2 - Core Feature Inoperative	Murugesan Kumar	Ajay	Pre-upgrade JAR doesn't able to run and getting exception. javax.security.sasl.SaslException: No common protection layer between client and server For above kerberos issue we've added "hadoop.rpc.protection=privacy" in hive-site.xml file - After that we faced the other for HDFS HA connection "java.lang.IllegalArgumentException: java.net ⁹⁵ .UnknownHostException:", For this issue we found /etc/hive/conf/conf.server was missing core-site.xml and hdfs-site.xml. Placed the needed files and the issue is resolved.	Pre-upgrade issues when HDP 2.6.5 to 3.1 in PERF cluster

⁹³ <https://support.hortonworks.com/s/case/50044000010sqi9AAA/detail>

⁹⁴ <https://support.hortonworks.com/s/case/50044000010slkCAAQ/detail>

⁹⁵ <http://java.net>

79	00209775(see page 3)	NULL is not inserting into Hive ACID Table through Spark	2/4/2019 8:54 AM	Closed	Support Analysis in Progress	Spark	S2-Core Feature Inoperative	Murugesan Kumar	Sarthak	Raised the case for NULL is not able to insert. Shared the required application logs to debug the same.	Application config issue. HWX fixed
78	00209772(see page 3)	Unable to read data from Hive in Spark through executeQuery	2/4/2019 8:45 AM	Closed	Support Analysis in Progress	Spark	S2-Core Feature Inoperative	Murugesan Kumar	David	Raised the case for read process from spark-submit command. Shared the required application logs to debug the same.	Similar case of 00208287
77	00209711⁹⁶	Post upgrade unable to start the YARN - TIMELINE SERVICE V2.0 READER	2/4/2019 4:55 AM	Closed	Support Analysis in Progress	YARN	S2-Core Feature Inoperative	Murugesan Kumar	Nilesh	post upgrade unable to bring up YARN Timeline reader service alone. Explained through webex about 30 min and shared the required logs. HWX will check and come back for this issue.	Upgrade issues when HD P 2.6.5 to 3.1

⁹⁶ <https://support.hortonworks.com/s/case/50044000010sdURAAY/detail>

76	00209280(see page 3)	Upgrade failure to move Hive tables from 2.6.5 to 3.1	1/31/2019 4:22 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Paablo	<p>Caused by: java.net⁹⁷.ConnectException: Error while authenticating with endpoint: http://cov3lhdpfst03.am.tsacorp.com:9292/kms/v1/?op=GETDELEGATIONTOKEN&renewer=rm%2Fcov3lhdpfst02.am.tsacorp.com%40IPA.AM.TSACORP.COM at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method) ~[?:1.8.0_141]</p> <p>Identified that, Ranger KMS/Admin was down. There are some ranger services installation/configuration issues. HWX reinstalled the required packages and aligned the configurations are properly.</p> <p>Post changes, Hive LLAP is brought up. Upgrade process are went smoothly.</p>	Upgrade issues when HD P 2.6.5 to 3.1
75	00209043(see page 3)	Upgrade failure to move Hive tables from 2.6.5 to 3.1	1/30/2019 2:48 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Paablo	<p>Next hive start services, we are blocked and getting exception: 2019-01-30T07:23:00,450 INFO [main]: metastore.RetryingMetaStoreClient (:()) - RetryingMetaStoreClient proxy=class org.apache.hadoop.hive.ql.metadata.SessionHiveMetaStoreClient ugi=hdfs-devcovhdp@IPA.AM.TSACORP.COM⁹⁸ (auth:KERBEROS) retries=24 delay=5 lifetime=0 2019-01-30T07:23:08,298 ERROR [main]: util.UpgradeTool (:()) - Bucketed table redi.csi_pends contains file hdfs://DEVNNHA/apps/hive/warehouse/redi.db/csi_pends/ penddatetime=2018-12-17</p> <p>There are some part files are causing this issue. Identified those and removed from hive DB.</p>	Upgrade issues when HD P 2.6.5 to 3.1

⁹⁷ <http://java.net/>

⁹⁸ <mailto:hdfs-devcovhdp@IPA.AM.TSACORP.COM>

74	00209012 ⁹⁹	Upgrade issues for Regenerate Missing Keytabs from HDP 2.6.5 to 3.1	1/29/2019 11:59 PM	Closed	Resolved by Customer	Ambari	S2 - Core Feature	Murugesan Kumar	-	<p>During upgrade process, faced issues to regenerate keytabs that, it created new keytab file as "yarn-ats.hbase*.keytab"</p> <p>Exception: resource_management.core.exceptions.Fail: User 'yarn-ats' doesn't exist</p> <p>Checked with internal IT team and created the user 'yarn-ats' and provided the appropriate permissions and the issue was resolved. asked HWX to close the case.</p>	Upgrade issues when HDP 2.6.5 to 3.1
73	00208952 ¹⁰⁰	Unable to upgrade from HDP 2.6.5 to 3.1	1/29/2019 1:33 PM	Closed	Resolved	Hive	S2 - Core Feature	Murugesan Kumar	Nitin Garg	<p>As part of upgrade, struck up the pre-upgrade checks for hive tables. Downloaded the JAR from HDP and tried to execute and got exception.</p> <p>1st Issue:</p> <p>Found Acid table: default.consumerwritertest 2019-01-29T16:38:13,171 ERROR [main] org.apache.hadoop.hive.upgrade.acid.PreUpgradeTool - PreUpgradeTool failed java.lang.IllegalArgumentException: java.net¹⁰¹.UnknownHostException: DEVNNHA To resolve this we added the hdfs-site and core-site.xml in /etc/hive/conf/conf.server/ location.</p> <p>2nd Issue:</p> <p>ERROR [main] org.apache.hadoop.hive.upgrade.acid.PreUpgradeTool - PreUpgradeTool failed java.io¹⁰².FileNotFoundException: File hdfs://DEVNNHA/apps/hive/warehouse/redi.db/trans_detail/oiddateyyyyymmdd=20171010</p>	Upgrade issues when HDP 2.6.5 to 3.1

⁹⁹ <https://support.hortonworks.com/s/case/50044000010sKDXAA2/detail>

¹⁰⁰ <https://support.hortonworks.com/s/case/50044000010sIR0AAM/detail>

¹⁰¹ <http://java.net>

¹⁰² <http://java.io>

										<p>does not exist. The data was not present on this table. So we drop the table.</p> <p>It went successful to proceed next step.</p>	
7 2	00 20 88 15 ¹⁰³	Unable to proceed upgrade from HDP 2.6.5 to 3.1 version	1/29/ 2019 5:47 AM	C l o s e d	Res olv ed	Am ba ri	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar	Br un o Ba ha m on de s(s ee pag e 3)	<p>As part of upgrade activity, cant't able to store the admin principal to proceed. Getting exception:</p> <p>After that when I do upgrade, it gives error "status: 403, message: Unable to sign in. Invalid username/password combination." Cant able to proceed further.</p> <p>Had long webex with HWX and we tried running the Kerberos service check, this is to test the credentials provided were correct, still hitting the same error mentioned before. There's a bug recently fixed on this, manual fix is to add a principal to the kdc with the format kadmin/KDC_FQDN, we did the following addprinc kadmin/ kerberos.ipa.am.tsacorp.com¹⁰⁴</p> <p>The next issue was that kerberos.ipa.am.tsacorp.com¹⁰⁵ was resolving to its canonical name, we added the following to /etc/hosts and upgrade started working.</p>	Upg rad e issu es whe n HDP 2.6. 5 to 3.1

¹⁰³ <https://support.hortonworks.com/s/case/50044000010sGFRAA2/detail>

¹⁰⁴ <http://kerberos.ipa.am.tsacorp.com/>

¹⁰⁵ <http://kerberos.ipa.am.tsacorp.com/>

71	00208527 ¹⁰⁶	Upgrade issues for Hive Metastore backup	1/27/2019 11:38 PM	Closed	Resolved by Customer	Hive	S2 - Core Feature	Minor	-	Issues faced to stop the Hive services in Ambari. So raised the case with HWX. With the logs identified, there are Ambari metrics are down to report to collector service. Due that unable to stop the Hive services. Post that, able to stop and taken the Hive metastore dump. Asked the HWX to close this.	Upgrade issues when HDP 2.6.5 to 3.1
70	00207966 ¹⁰⁷	In Structured streaming, to_json on array<struct> column is supposed to make it a single json string but while saving into hive, the data is getting split the attributes value and saving those into other multiple columns.	1/23/2019 6:24 AM	Closed	Work in progress with HWX	Spark	S2 - Core Feature	Minor	Akash Shinde (see page 3)	Shared the required inputs with HWX and the debugging is still going on.	Config issue. Adjusted in query as per link ¹⁰⁸

¹⁰⁶ <https://support.hortonworks.com/s/case/50044000010sBVWAA2/detail>

¹⁰⁷ <https://support.hortonworks.com/s/case/500440000108bSHAAY/detail>

¹⁰⁸ https://docs.hortonworks.com/HDPDocuments/HDP3/HDP-3.1.0/integrating-hive/content/hive_use_hwc_for_streaming.html

69	00207953 ¹⁰⁹	In Spark Structured Streaming Column containing timestamp value is not stored with proper values in hive table with column type timestamp or string	1/23/2019 5:24 AM	Closed	Work in progress with HWX	Spark	S2 - Core Feature Initiative	Murugesan Kumar	Harigopal	Shared the required inputs with HWX and the debugging is still going on.	BUG accepted by HWX and provided the hotfix
68	00207946 ¹¹⁰	Hive Metastore is DOWN	1/23/2019 4:59 AM	Closed	Resolved	Hive	S2 - Core Feature Initiative	Vaibhav	Vaibhav	In DEV cluster Masternode 02 was facing issues due to /var consumed 100% . Post cleared the old logs unable to bring up the Hive metastore services. Shared the logs with HWX and identified the existing service hasnt been stopped properly. Killed the service forcefully and brought up the services back.	Issues for the old process ID hasnt been killed properly.

¹⁰⁹ <https://support.hortonworks.com/s/case/500440000108bHnAAI/detail>

¹¹⁰ <https://support.hortonworks.com/s/case/500440000108bEKAAY/detail>

67	00207943 ¹¹¹	Hive server2 and Metastore is DOWN	1/23/2019 4:28 AM	Closed	Resolved	Hive	Spark	Surgesan Kumar	Vai bhav	<p>In DEV cluster Masternode 02 was facing issues due to /var consumed 100% . Post cleared the old logs unable to bring up the Hive server2.</p> <p>Shared the logs with HWX and identified the existing service hasnt been stopped properly. Killed the service forcefully and brought up the services back.</p>	Issues for the old process ID hasnt been killed properly.
66	00207913 (see page 3)	Count operation on a spark dataframe is not giving correct result	1/23/2019 1:25 AM	Closed	Work in progress with HWX	Spark	Spark	Surgesan Kumar	Ranjith Pullu (see page 3)	Shared the required inputs with HWX and the debugging is still going on.	Documentation issues from HWX. Spark config issue.

¹¹¹ <https://support.hortonworks.com/s/case/500440000108b8gAAA/detail>

65	00207873 ¹¹²	Unable to use hivesession .executeQuery HWC API.	1/22/2019 7:16 PM	Closed	Duplicate case created by HWX	Spark	S2 - Core Feature Incorporated	Ranjith Pulluru (see page 3)	Ranjith Pulluru (see page 3)	For read issues already we reported 00207124 ¹¹³ case. Ranjith was created case by mistake	Duplicate case created by HWX
64	00207636 ¹¹⁴	HDP Upgrade request for DEV cluster	1/21/2019 9:51 AM	Closed	Opened	Ambari/HDP	Admin	Murugesan Kumar	Mayank Gupta (see page 3)	Uploaded the required details about the cluster	DEV cluster upgrade query from HDP 2.6.5 to HDP 3.1

¹¹² <https://support.hortonworks.com/s/case/500440000108aBMAAY/detail>

¹¹³ <https://support.hortonworks.com/s/case/500440000108Fi5AAE/detail>

¹¹⁴ <https://support.hortonworks.com/s/case/500440000108UmlAAE/detail>

63	00207533 ¹¹⁵	Clarifications on Hive ORC Table properties while creating Hive default Transactional table	1/20/2019 11:19 PM	Closed	Opened the case	Hive server2	Spark - Core Feature Inoperative	Murugesan Kumar	Rohit Rami Mahotra	Clarifications on Hive ORC Table properties while creating Hive default Transactional table for the below queries: 1. what default compression is used in Default transactional table 2. default delimiter in Hive default transactional table 3. all these properties are used in Hive 2 similar properties are required for Hive 3 also to create hive transactional table to get more performance	Clarification from HWX query
62	00207534 ¹¹⁶	HiveWarehouse connector not able to write data into complex data types in Hive acid tables	1/20/2019 11:25 PM	Closed	Create the case with the inputs	Spark SQL	Spark - Core Feature Inoperative	Murugesan Kumar	Ranjith Pulluru (see page 3)	It is been observed that Hivewarehouse connector which is used in SparkStructured streaming is not able to write records into hive acid table columns which have complex data types likes ARRAY<STRUCT<>>. It stores only NULL values but seems like it is able to recognize the attributes inside.	BUG accepted by HWX

¹¹⁵ <https://support.hortonworks.com/s/case/500440000108TTVAA2/detail>

¹¹⁶ <https://support.hortonworks.com/s/case/500440000108TTaAM/detail>

6 1	00207528 ¹¹⁷	Not able to read Complex datatype (Array value) from spark through HWC	1/20/2019 10:41 PM	Closed	Opened the case	Spark SQL	S2 - Core Feature Inoperative	Murugesan Kumar	Ranjith	Opened the case with the exception reported. HWX confirms it's a bug in spark. so engaged Spark SME on this case.	BUG accepted from HWX
--------	---	--	--------------------	--------	-----------------	-----------	-------------------------------	-----------------	---------	--	-----------------------

¹¹⁷ <https://support.hortonworks.com/s/case/500440000108TPdAAM/detail>

60	00207328 ¹¹⁸	Write into Hive transactional table in hive 3 by using HWC - Cluster Mode	1/18/2019 1:19 AM	Closed	Resolved	Spark SQL	Spark - Core Feature Incorporative	Murugesan Kumar	Sarthak Saxena	<p>Getting exception for cluster mode and uploaded the application logs to investigate.</p> <p>Had webex and found that,</p> <ul style="list-style-type: none"> - The application was trying to make a PLAIN connection to zookeeper due to which it was failing. We added the principal in the JDBC URL which resolved the issue. - Later the application failed with a Kerberos issue. We added --keytab and --principal in the spark-submit which resolved the issue(added at /etc/spark2/conf/spark2-defaults.conf). - We also observed that Hive LLAP was in accepted state. We restarted the daemon and post that it was working properly. <p>Added the property at spark.sql.hive.hiveserver2.jdbc.url under custom spark2-default.conf and now able to execute through cluster mode.</p> <p>Default hive principal doesnt support at HDP 3.1. Explicitly we need to include that.</p>	<p>Documentation issue.</p> <p>As per the guide, principal configured as separate parameter. But when we add in the URL only we are able to fix this issue.</p>
----	-------------------------	---	-------------------	--------	----------	-----------	------------------------------------	-----------------	----------------	---	---

¹¹⁸ <https://support.hortonworks.com/s/case/500440000108JmGAU/detail>

59	00207322 ¹¹⁹	Clarification on Hive Transactional tables creation in Hive3	1/18/2019 1:01 AM	Completed	Work in progress with HWX	Spark SQL	Scorpio Feature Inoperative	Murugesan Kumar	Nisha Vaidya (see page 3)	Clarification required to create Hive Transactional tables in Hive3. HDP 3.1 supports partially Primary key constraints. HWX would verify further and come back with the solution.	Clarification from HWX query
58	00207312 ¹²⁰	ACID Table Update/Delete is not working through HWC	1/18/2019 12:00 AM	Completed	Work in progress with HWX	Spark SQL Hive server2	Scorpio Feature Inoperative	Murugesan Kumar	Ranjith Puthuttu (see page 3) Nidhin Divakaran	<p>hanging the job and YARN UI application is in progress state. gave enough resources.</p> <p>Webex debugging in progress with HWX. Suggestion gave us problem with Hive update/delete.</p> <p>So need to create a separate ticket under Hive services. Post fix will come back to this case for Spark.</p> <p>Changed the ticket to Hive services instead of Spark.</p> <p>resource crunch on the cluster. we installed only 2 node managers(05/06 vms). So need to add one more nodemanager service in the 04 vm. waiting for IT team to install the same.</p> <p>Installed the NM at 04 vm. But still the application hangs and reported to HWX. It was hanging due to insufficient resources. Increased the queue percentage and now we able to execute through client/cluster modes.</p>	POC Configuration issues.

119 <https://support.hortonworks.com/s/case/500440000108JkFAAU/detail>

120 <https://support.hortonworks.com/s/case/500440000108JfFAAU/detail>

57	00207141 ¹²¹	Current timestamp function is return null when writing into hive in streaming	1/17/2019 1:29 AM	Cloned	Work in progress with HWX	Spark Streaming	S2 - Core Feature Initiative	Murugesan Kumar	Sarthak Saxena	Current timestamp function is return null when writing into hive in streaming with the following spark-submit command Work in progress with HWX team. HWX accepted as BUG and the Engineering team engaged for fixing.	BUG accepted by HWX
56	00207124 ¹²²	Read into Hive transactional table in hive 3 by using HWC	1/16/2019 11:10 PM	Cloned	HWX needs time to check with their team	Spark	S2 - Core Feature Initiative	Murugesan Kumar	Sarthak Saxena	getting exception while read into Hive transactional table by HWC. 19/01/17 00:38:09 ERROR Executor: Exception in task 0.1 in stage 0.0 (TID 1) java.lang.RuntimeException: java.io ¹²³ .IOException: No service instances found in registry Identified that, issues from Hive LLAP side instead of Spark. Shared the interactive.log file to analyse with their Hive team. The issues with spark-default.conf file. Changed and restarted the services.	Configuration issues in spark configuration.

¹²¹ <https://support.hortonworks.com/s/case/500440000108FrgAAE/detail>

¹²² <https://support.hortonworks.com/s/case/500440000108Fi5AAE/detail>

¹²³ <http://java.io>

55	00206925 ¹²⁴	Zookeeper acl create failure for Hive user	1/16/2019 12:21 AM	Closed	Resolved	Zookeeper	S2 - Core Feature Implementation	Murugesan Kumar	Magdum Sajjan	<p>Ambari -> Services -> ZooKeeper -> Configs -> Advanced zookeeper-env and add below attribute to zookeeper-env template.</p> <pre>export SERVER_JVMFLAGS="\$SERVER_JVMFLAGS -Dzookeeper.skipACL=yes -Djava.security.auth.login.config={{zk_server_jaas_file}}"</pre> <p>[zk: hdp243v:2181(CONNECTED) 0] getAcl /llap-sasl/user-poc_hive <<==== Save this output [zk: hdp243v:2181(CONNECTED) 1] setAcl /llap-sasl/user-poc_hive world:anyone:r:sasl:hive:sdrwa</p> <p>Post creation of ACL now revert back the changes in Zookeeper. Now works fine.</p>	Issues with poc_hive user as documentation.
54	00206742 ¹²⁵	Count(*)/Count(1) or count(column) not giving proper results	1/14/2019 11:03 PM	Closed	Resolved	Hiveserver2	S2 - Core Feature Implementation	Murugesan Kumar	Nitin Garg	<p>see on web-ex that all the jobs were in accepted state and getting below message in the RM UI. Application is added to the scheduler and is not yet activated. User's AM resource limit exceeded. Details : AM Partition = <DEFAULT_PARTITION>; AM Resource Request = <memory:2048, vCores:1>; Queue Resource Limit for AM = <memory:6144, vCores:1>; User AM Resource Limit of the queue = <memory:2048, vCores:1>; Queue AM Resource Usage = <memory:1024, vCores:1>;</p> <p>So we changed the queue capacity and refresh the yarn queues. After this we were able to run the query.</p>	Configuration issue in POC Cluster.

¹²⁴ <https://support.hortonworks.com/s/case/500440000108Bh4AAE/detail>

¹²⁵ <https://support.hortonworks.com/s/case/5004400001086gOAAQ/detail>

53	00206606 ¹²⁶	Read/Write into Hive transactional table in hive 3 by using HWC	1/14/2019 7:00 AM	Closed	Service Recovered - Pending from SME	Spark	S2 - Core Feature Inoperative	Murugesan Kumar	Harikumar	Multiple Webex sessions taken and identified as bug from HDP product. 'Hive' user hardcoded to do authorize right access. When we change into hive-site.xml and hive interactive-site.xml as poc_hive from hive user, able to insert data into hive from Spark.	This is incorrect documentation from Cloudera. Actually we use common LDAP server for all the 3 PD environments. Already the name prefix suggested by HWX.
----	-------------------------	---	-------------------	--------	--------------------------------------	-------	-------------------------------	-----------------	-----------	---	--

¹²⁶ <https://support.hortonworks.com/s/case/500440000108306AAI/detail>

52	00206522 ¹²⁷	Upgrade to HDP 3.1 from HDP 2.6.5	1/13/2019 9:14 PM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar	Akshay Jadhav (see page 3)	HWX suggested to open a new request for Upgradation as separate action plan which requires downtime.	Query case before doing the upgrade
51	00204327 ¹²⁸	Require pom.xml for HDP 3.1.0 and sandbox for 3.1.0	12/21/2018 7:43 AM	Closed	Resolved	Sandbox	S2 - Core Feature Inoperative	Murugesan Kumar	Javier Fuentes	HWX provided the pom.xml file but sandbox 3.0 was not available to share.	

¹²⁷ <https://support.hortonworks.com/s/case/5004400001081RyAAI/detail>

¹²⁸ <https://support.hortonworks.com/s/case/500440000106sFNAAY/detail>

50	202277 ¹²⁹	Exception while concurrent reading/writing from Hive ACID tables using Spark: Caused by: java.util.concurrent.ExecutionException: java.lang.NumberFormatException: For input string: "0004880_000"	12-11-2018 05:58	Closed	Enhancement Request Logged	Spark	Spark 2 - Core Feature Inoperative	Aleksi Filushin ¹³⁰	Nidhin Divakaran	This issue cant be fixed at 2.6.5 version	
49	200945 ¹³¹	Issues with Hivemetastore - Postgresql	12-03-2018 21:09	Closed	Resolved	Hive	Spark 2 - Core Feature Inoperative	Murugesan Kumar ¹³²	Kartik Nadar	Issues with Hive metastore - Postgres. Fixed with the workaround.	

¹²⁹ <https://support.hortonworks.com/s/case/5004400001067X0AAI/detail>

¹³⁰ <https://support.hortonworks.com/s/detail/0034400001sFtcAAC>

¹³¹ <https://support.hortonworks.com/s/case/5004400000yKXFTAA4/detail>

¹³² <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

4 8	20 08 05 ¹³³	saving into Hive ACID table taking time from Hive external table	12-03 -2018 08:15	C l o s e d	Pen din g Cus tom er Inp ut	Hiv e	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar ¹³⁴	Nit in Ga rg	Updated with required logs to reduce the Hive external table loading time	
4 7	20 08 03 ¹³⁵	Saving into HDFS as ORC format taking time	12-03 -2018 08:12	C l o s e d	Pen din g Cus tom er Inp ut	Sp ark	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar ¹³⁶	Ra nji th Pu llu ru	Code fixing is going on to check the response time	

¹³³ <https://support.hortonworks.com/s/case/5004400000yKTHnAAO/detail>

¹³⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹³⁵ <https://support.hortonworks.com/s/case/5004400000yKTGTAA4/detail>

¹³⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

4 6	200774 ¹³⁷	Can't connect from Spark to Hive LLAP	12-03-2018 06:17	Closed	Support Analysis in Progress	Spark	Spark2 - Core Feature Inoperative	Aleksi Filushin ¹³⁸	David Bompart Marquez	<p>We have discussed this with engineering, and it has been confirmed that this is a known requested feature that has been a community marked as "Won't fix", for your reference there are two Apache JIRAs related to this:</p> <p>- https://issues.apache.org/jira/browse/SPARK-22016 - https://issues.apache.org/jira/browse/SPARK-21063</p>	
4 5	199873 ¹³⁹	Spark2 uses hive-metastore-1.21 at HDP 2.6.5	11/27/2018 1:06 AM	Closed	Resolved	Spark	Spark2 - Core Feature Inoperative	Murugesan Kumar ¹⁴⁰	Sarthak Saxena	<p>There are jars misplaced in the spark-library. Removed the jars which are related to 2.1 metastore. Now thrift server is UP.</p>	

¹³⁷ <https://support.hortonworks.com/s/case/5004400000yKShmAAG/detail>

¹³⁸ <https://support.hortonworks.com/s/detail/0034400001sFtcAAC>

¹³⁹ <https://support.hortonworks.com/s/case/5004400000yJzmmAAC/detail>

¹⁴⁰ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

4 4	19 90 59 ¹⁴¹	Hive LLAP Join query takes 30 sec to be reduced	11/20/2018 4:26 AM	Closed	Support Analysis in Progress	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁴²	Paolo Idiazquez	Updated with required logs to reduce the Hive LLAP Reponse. HWX works with their SME	
4 3	19 87 86 ¹⁴³	Configuration issue on Spark streaming job	11/18/2018 9:49 PM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁴⁴	Manuel Calvo Rodriguez	we found an issue related to keystore.jks not found so we asked to create the path in each worker node or run the job using --files ./pathToFiles - this fixed the issue	

¹⁴¹ <https://support.hortonworks.com/s/case/5004400000yJbOxAk/detail>

¹⁴² <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁴³ <https://support.hortonworks.com/s/case/5004400000yJTw8AAG/detail>

¹⁴⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

4 2	19 84 61 ¹⁴⁵	Hive insertion issue on the TRE feature flow	11/15/2018 8:17 AM	Closed	Pending Customer Input	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁴⁶	David Bompartz	As per David's recommendations team changed the code. Now we are in testing phase for further updates.	
4 1	19 84 59 ¹⁴⁷	Hive insertion issue on the Main exec flow	11/15/2018 8:12 AM	Closed	Support Analysis in Progress	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁴⁸	Matt Barr	As per David's recommendations team changed the code. Now we are in testing phase for further updates.	

¹⁴⁵ <https://support.hortonworks.com/s/case/5004400000yJHJZAA4/detail>

¹⁴⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁴⁷ <https://support.hortonworks.com/s/case/5004400000yJHGnAAO/detail>

¹⁴⁸ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

40	198448 ¹⁴⁹	Supporting tools for deployment	11/15/2018 7:26 AM	Closed	No Customer Response	Ambari	S4 - Request for Information	MiddhounPoli ¹⁵⁰	Shivraj Singh Rathore	Recommended tools	
39	197849 ¹⁵¹	Requesting resolution plan for known Security Vulnerabilities with HDP 2.6.5 version.	11-12-2018 17:54	Closed	Pending Engineering Input	Spark	S4 - Request for Information	Kesarvande ¹⁵²	Sarthak Saxena	Patching in progress from SME team	

¹⁴⁹ <https://support.hortonworks.com/s/case/5004400000yJGvPAAW/detail>

¹⁵⁰ <https://support.hortonworks.com/s/detail/0034400001oJsY8AAK>

¹⁵¹ <https://support.hortonworks.com/s/case/5004400000yJ2gqAAC/detail>

¹⁵² <https://support.hortonworks.com/s/detail/0034400001vq1jYAAQ>

38	196033 ¹⁵³	cache or persist takes long time in spark-submit jobs	10/31/2018 2:14 AM	Closed	Resolved by Customer	Spark	S2 - Core Feature	Murugesan Kumar ¹⁵⁴	Sarthak Saxena	Code fixing resolved the cache issues	
37	196024 ¹⁵⁵	Submitting Spark job through Livy takes more time, than it was before	10/31/2018 1:35 AM	Closed	No Customer Response	Spark	S2 - Core Feature	Aleksi Filushin ¹⁵⁶	Pankaj Kadam		

¹⁵³ <https://support.hortonworks.com/s/case/5004400000yl4hgAAC/detail>

¹⁵⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁵⁵ <https://support.hortonworks.com/s/case/5004400000yl4dUAAS/detail>

¹⁵⁶ <https://support.hortonworks.com/s/detail/0034400001sFtciAAC>

3 6	19 51 10 ¹⁵⁷	Spark submit fails with class def not found and confluent class not found issue related	10/24 /2018 10:46 PM	C l o s e d	Res olv ed	Sp ark	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar ¹⁵⁸	Ak sh ay Ma nk u m ba re	resolved this issue by passing required confluent jars in spark submit command.	
3 5	19 49 23 ¹⁵⁹	Unable to access kafka cluster from Spark-submit using srvredi service account	10/24 /2018 2:50 AM	C l o s e d	Res olv ed	Kaf ka	S 2 - C or e Fe at ur e In o p er at iv e	M ur ug es an K u m ar ¹⁶⁰	Ak sh ay Ma nk u m ba re	Created ACLs for srvredi user account and now able to construct Kafka consumer through srvredi user account.	

157 <https://support.hortonworks.com/s/case/5004400000yHdmLAAS/detail>

158 <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

159 <https://support.hortonworks.com/s/case/5004400000yHYG8AAO/detail>

160 <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

3 4	19 49 03 ¹⁶¹	Inserting static data into hive table with specific schema access the default database	10/24 /2018 12:15 AM	Closed	No Customer Response	Hive	Spark - Core Feature Inoperative	Murugesan Kumar ¹⁶²	Nisha Vaidya	haven't stored now in default DB. Waiting for the next attempt if in case occurred	
3 3	19 47 19 ¹⁶³	Livy connectivity issues	10/23 /2018 3:46 AM	Closed	Resolved	Spark	Spark - Core Feature Inoperative	Aleksi Filushin ¹⁶⁴	Harigopal Matta	always pass the hive-site.xml that's present in /etc/spark/conf or /etc/spark2/conf(basing on spark1/spark2 job) . If you change anything in hive-site.xml of these then you should manually put this file back to the hdfs location from where you are trying to use.	

¹⁶¹ <https://support.hortonworks.com/s/case/5004400000yHXBAAW/detail>

¹⁶² <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁶³ <https://support.hortonworks.com/s/case/5004400000yHSKPAA4/detail>

¹⁶⁴ <https://support.hortonworks.com/s/detail/0034400001sFtciAAC>

3 2	19 46 83 ¹⁶⁵	Unable to run hive join query in beeline getting vertex error	10/22/2018 11:15 PM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁶⁶	Karthik Gopal	The issue happens when the hdfs delegation token is not retained for a longer time as needed by the hive user to access hdfs and we get the error as reported in your issue. include custom tez-site.xml with tez.cancel.delegation.tokens.on.completion=false and requires hive/Oozie to restart	
3 1	19 32 63 ¹⁶⁷	Storing to Hive tables Takes a long time	10-12-2018 11:24	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁶⁸	Hari Gopal Matta	Code levels changes are in progress as per David's recommendations. If in case issue persists HWX suggested to open a new case	

¹⁶⁵ <https://support.hortonworks.com/s/case/5004400000yHRj4AAG/detail>

¹⁶⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁶⁷ <https://support.hortonworks.com/s/case/5004400000y3TDpAAM/detail>

¹⁶⁸ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

30	192672 ¹⁶⁹	Spark jobs monitoring - need help	10-09-2018 13:21	Closed	No Customer Response	Spark	S3 - Minor Feature Inquiry ¹⁷⁰	Mihun Polisetty ¹⁷⁰	Xin Wang	Recommended tools	
29	192639 ¹⁷¹	Hive LLAP Join query takes 30 sec to be reduced	10-09-2018 10:47	Closed	Resolved	Hive	S2 - Core Feature Inquiry ¹⁷²	Murugesan Kumar ¹⁷²	Paolo Jung	After multiple webex sessions, shared the logs to debug further. Case transferred to #199059 due to vacation	

¹⁶⁹ <https://support.hortonworks.com/s/case/5004400000y1oSdAAI/detail>

¹⁷⁰ <https://support.hortonworks.com/s/detail/0034400001oJsY8AAK>

¹⁷¹ <https://support.hortonworks.com/s/case/5004400000y1iIVAAY/detail>

¹⁷² <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

28	192375 ¹⁷³	Suitable driver not found error in Spark-submit transflow job	10-08-2018 06:13	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁷⁴	Hari Gopal Matta	<p>The problem is not with the jdbc jar not being available in classpath, as its available in both driver and executor class path, As found in webex, jdbc driver class name option is not set in spark.read.format(jdbc) api, for oracle this is compulsory.</p> <p>After setting the jdbc driver option ,the job is running successfully.</p>
27	192355 ¹⁷⁵	Unable to find jaas.conf file while running ingestion job in YARN mode	10-08-2018 04:15	Closed	Resolved	Kafka	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁷⁶	Nilesh Gangaram Parmar	Issues with srvredi access. IT team created principal for srvredi to access kafka cluster.

¹⁷³ <https://support.hortonworks.com/s/case/5004400000y0ywtAAA/detail>

¹⁷⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁷⁵ <https://support.hortonworks.com/s/case/5004400000y0wPnAAI/detail>

¹⁷⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

26	192198 ¹⁷⁷	Externalizing Configurations Resulting in an Error	10-05-2018 09:25	Closed	No Customer Response	Spark	S2 - Core Feature Inoperative	Kesa v An de ¹⁷⁸	Daniel Ruiz Robo	
25	191571 ¹⁷⁹	Resource Manager HA not functioning properly - Both RMs going to standby/standby state, split-brain scenario. Other clusters with same config, works fine.	10-02-2018 06:29	Closed	Resolved	YARN	S2 - Core Feature Inoperative	Mihun P olise tt y ¹⁸⁰	Ramakrishna Reddy Kanchi	Issues with Kerberos ticket and fixed as per the guidance from HWX

¹⁷⁷ <https://support.hortonworks.com/s/case/5004400000xn0OVAAY/detail>

¹⁷⁸ <https://support.hortonworks.com/s/detail/0034400001vq1jYAAQ>

¹⁷⁹ <https://support.hortonworks.com/s/case/5004400000xkjReAAI/detail>

¹⁸⁰ <https://support.hortonworks.com/s/detail/0034400001oJsY8AAK>

2 4	19 10 56 ¹⁸¹	Spark broadcast issue	9/27/2018 4:01 PM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Kesa Vande ¹⁸²	Ravi Bhargava	We could see OOME in the spark driver logs. HWX suggested to increase the heap by 20% to 28GB for both driver and executor. We could see the job went past that issue however failed for some other transformation issue.	
2 3	19 08 30 ¹⁸³	Hive query - select count(*) from table doesnt work	9/26/2018 11:53 PM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁸⁴	Nitin Garg	Running the job with user srvredi. And this users doesn't have permission on the all the Node managers. Due to which your job was failing. After changing the permission you were able to run the query.	

¹⁸¹ <https://support.hortonworks.com/s/case/5004400000xJkZAAU/detail>

¹⁸² <https://support.hortonworks.com/s/detail/0034400001vq1jYAAQ>

¹⁸³ <https://support.hortonworks.com/s/case/5004400000xO2dEAAS/detail>

¹⁸⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

2 2	190189 ¹⁸⁵	Spark Execution gets failed	9/23/2018 11:57 PM	Closed	No Customer Response	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁸⁶	Harigopal Matta	HWX suggested to do the code fixes: Broadcast joins are different and its based on spark deciding when to do such broadcast of dataframe basing on parameter spark.sql.autoBroadcastJoinThreshold (which is 10MB by default). Also, you shouldn't try to do broadcast join on larger datasets just by increasing the spark.sql.autoBroadcastJoinThreshold	
2 1	189977 ¹⁸⁷	While writing data to Hive sink during our Spark structure streaming job into 2.3.0 we get the following exception.	9/21/2018 1:42 AM	Closed	No Customer Response	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁸⁸	Ranjith Pulluru	Code issues	

185 <https://support.hortonworks.com/s/case/5004400000xMvMnAAK/detail>

186 <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

187 <https://support.hortonworks.com/s/case/5004400000xMHRCAAW/detail>

188 <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

20	189307 ¹⁸⁹	Unable to do spark-submit getting No driver found exception	9/18/2018 3:50 AM	Closed	Resolved	Spark	Spark2 - Core Feature Inoperative	Murugesan Kumar ¹⁹⁰	Sarthak Saxena	HWX fixed the issue by placing the oracle jar to "/usr/hdp/current/spark2-client/jars/"	
19	189203 ¹⁹¹	Ambari-infra-solr on one out of three instances dies with OOM error	9/17/2018 12:35 PM	Closed	Resolved	Ambari	Spark3 - Minor Feature Inoperative	Mihun Pottiy ¹⁹²	Alejandro Arriola	HWX gave the workaround: https://issues.apache.org/jira/browse/AMBARI-22684 more information about this bug can be found here: https://risdenk.github.io/2017/12/18/ambari-infra-solr-ranger.html	

¹⁸⁹ <https://support.hortonworks.com/s/case/5004400000xL6XxAAK/detail>

¹⁹⁰ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁹¹ <https://support.hortonworks.com/s/case/5004400000xKuYSAA0/detail>

¹⁹² <https://support.hortonworks.com/s/detail/0034400001oJsY8AAK>

18	189084 ¹⁹³	Facing issues after upgrading to HDP 2.6.5(Spark 2.3.0)	9/17/2018 4:01 AM	Closed	Duplicate	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁹⁴	Ranjith Pulluru	HWX resolved with some parameters tuning in the cluster. As well Duplicate of this case	
17	189005 ¹⁹⁵	Unable to start some services in HDP 2.6.4	9/15/2018 12:03 AM	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Murugesan Kumar ¹⁹⁶	Anchal Sharma	Issues at ACI side We were trying to upgrade the cluster and found issues with our Satellite repository server. Our IT team has fixed the issues and upgraded the cluster too.	

¹⁹³ <https://support.hortonworks.com/s/case/5004400000xKgbKAAS/detail>

¹⁹⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

¹⁹⁵ <https://support.hortonworks.com/s/case/5004400000xKKYuAAO/detail>

¹⁹⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

1 6	188984 ¹⁹⁷	Need help to resolve security configuration issue - Service account not able to access encryption zone.	9/14/2018 2:09 PM	Closed	No Customer Response	Spark	S2 - Core Feature Inoperative	Kevin ¹⁹⁸	Xin Wang	The property spark.authenticate.secret is mainly used to force using a secret key. As you know, you are hitted https://issues.apache.org/jira/browse/SPARK-24062 . And the property if to workaround the BUG.	
1 5	188455 ¹⁹⁹	HDP 2.6.5.0 Upgrade issue - unable to register version using VDF file	09-11-2018 14:53	Closed	Resolved	Ambari	S2 - Core Feature Inoperative	Trent Melcher ²⁰⁰	Akash Singh	HWX provided the steps to fix the upgrade issues	

¹⁹⁷ <https://support.hortonworks.com/s/case/5004400000xKFRiAAO/detail>

¹⁹⁸ <https://support.hortonworks.com/s/detail/0034400001vq1jYAAQ>

¹⁹⁹ <https://support.hortonworks.com/s/case/5004400000x2MTBAA2/detail>

²⁰⁰ <https://support.hortonworks.com/s/detail/0034400001vq1fqAAA>

1 4	188417 ²⁰¹	Hive LLAP performance is very slow, not meeting business SLAs.	09-11-2018 12:00	Closed	Resolved	Hive	Spark - Core Feature Inoperative	Kesaavan ²⁰²	Pablo Idiazquez	Post upgradation from 2.6.4 to 2.6.5 postgres metastore altered. ALTER TABLE "PART_COL_STATS" ADD COLUMN "BIT_VECTOR" BYTEA; ALTER TABLE "TAB_COL_STATS" ADD COLUMN "BIT_VECTOR" BYTEA;	
1 3	187457 ²⁰³	Spark 2.3.0 not able to add with column and notable to write into hive by using hive sink	09-05-2018 05:28	Closed	No Customer Response	Spark	Spark - Core Feature Inoperative	Murugesan Kumar ²⁰⁴	Ranjith Pulluru	Security issues with spark-submit and fixed the same.	

201 <https://support.hortonworks.com/s/case/5004400000x2Hy1AAE/detail>

202 <https://support.hortonworks.com/s/detail/0034400001vq1jYAAQ>

203 <https://support.hortonworks.com/s/case/5004400000x0BDaAAM/detail>

204 <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

1 2	187219 ²⁰⁵	Unable to connect Hive from spark-summit	09-04-2018 02:00	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ²⁰⁶	Ranjith Pulluru	spark maven dependency added in the pom.xml file	
1 1	187043 ²⁰⁷	Unable to connect Hive from spark-summit	09-01-2018 02:13	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ²⁰⁸	Hari Gopal Matta	Issues with pom.xml and fixed	

²⁰⁵ <https://support.hortonworks.com/s/case/5004400000wzljRAAI/detail>

²⁰⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

²⁰⁷ <https://support.hortonworks.com/s/case/5004400000wzJfoAAE/detail>

²⁰⁸ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

10	186855 ²⁰⁹	Spark2 History Service - Enable HTTPS issue	8/30/2018 1:22 PM	Closed	Resolved	Spark	S3 - Minor Feature Improvement	Trent Miller ²¹⁰	Andrew Miller	<p>We found there is an error / mistake in the Spark documentation where spark.ssl.enabledAlgorithms should not be set, as Spark will negotiate the most secure and compatible cipher automatically. So we removed it from spark2-defaults and restarted Spark History Server and it works now to load the SHS URL.</p> <p>The internal JIRA BUG-109784 was opened August 23, 2018 to get this error corrected in the Spark documentation.</p> <p>BUG-109784 https://hortonworks.jira.com/browse/BUG-109784</p>	
9	186333 ²¹¹	Spark doesn't see Hive table when it's run through Livy	8/28/2018 5:44 AM	Closed	Resolved	Spark	S2 - Core Feature Improvement	Aleksi Filushin ²¹²	Harri Gopal Matta	<p>hive-site.xml thats passed using /etc/spark2/conf directory is a local directory not a hdfs one, hence the problem.</p> <p>might do this also, from /etc/spark2/conf, copy the hive-site.xml to any hdfs directory and pass it in your json file as follows: example: hadoop fs -put /etc/spark2/conf/hive-site.xml /user/spark/hive-site.xml</p>	

²⁰⁹ <https://support.hortonworks.com/s/case/5004400000wyyDMAAY/detail>

²¹⁰ <https://support.hortonworks.com/s/detail/0034400001vq1fqAAA>

²¹¹ <https://support.hortonworks.com/s/case/5004400000wldKkAAK/detail>

²¹² <https://support.hortonworks.com/s/detail/0034400001sFtciAAC>

8	186129 ²¹³	While trying to read HIVE transactional table from Spark 2.2.1, we are not able to do it	8/27/2018 6:14 AM	Closed	Resolved	Spark	S2 - Core Feature Inoperative	Murugesan Kumar ²¹⁴	Harish Gopal Matta	Issues with 2.6.4 and HWX suggested to do upgrade to 2.6.5	
7	186079 ²¹⁵	Couldnt access the Hive LLAP Beeline from normal user's principal	8/26/2018 11:30 PM	Closed	Resolved	Hive	S3 - Minor Feature Inoperative	Murugesan Kumar ²¹⁶	Nitin Garg	to use hive principal "hive/_HOST@HDP.LOCAL" instead of user's (srvredi) principal	

²¹³ <https://support.hortonworks.com/s/case/5004400000wUmhAAG/detail>

²¹⁴ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

²¹⁵ <https://support.hortonworks.com/s/case/5004400000wITOpAAO/detail>

²¹⁶ <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

6	185503 ²¹⁷	Need guidance to perform some performance testing/benchmarking for Hive LLAP	8/22/2018 6:48 AM	Closed	Referred to Consulting/Sales	Hive	S2 - Core Feature Implementation	Midhun Perera De Mattos	Erna	<p>You can use the following tool: https://github.com/hortonworks/hive-testbench/tree/hdp26</p> <p>Download and follow the install instructions. Possible you need a consultancy, then I would recommend Hortonwork's Professional Services. Please contact us at 1.855.0.HORTON ref: https://hortonworks.com/services/professional/#implementation</p> <p>HWX suggested to take professional services to hands on this Hive LLAP</p>
5	185339 ²¹⁹	Knox - LDAP encrypt password issue	8/21/2018 10:42 AM	Closed	Resolved	Knox	S3 - Minor Feature Implementation	Trent Melcher	Hemanth Sancha	<p>KnoxCli is just a testing tool and we have bug raised for the same KNOX-745 is not resolved yet. Security SME Vipin is working on it.</p> <p>Yes, the authentication will work, it is just bug you are hitting on. https://issues.apache.org/jira/browse/KNOX-745</p> <p>Trent confirmed as KnoxCLI is working as expected.</p>

²¹⁷ <https://support.hortonworks.com/s/case/5004400000w1l3AAC/detail>

²¹⁸ <https://support.hortonworks.com/s/detail/0034400001oJsY8AAK>

²¹⁹ <https://support.hortonworks.com/s/case/5004400000wHwn3AAC/detail>

²²⁰ <https://support.hortonworks.com/s/detail/0034400001vq1fqAAA>

4	185284 ²²¹	Hive LLAP Performance Issue	8/21/2018 7:48 AM	Closed	Duplicate	Hive	S2 - Core Feature Inoperative	Murugesan Kumar ²²²	Juan Chietta	There is a case (# 00185198) already opened. By mistake the new case has been opened.	
3	185203 ²²³	Add a new HDP Support portal User	8/20/2018 11:19 PM	Closed	Resolved	Administration (non tech)	S4 - Request useres for Information	Hariprasad Allak ²²⁴	Rishabh Shukla	Created a new account for Murugesan Kumar and closed the account for Hariprasad	

²²¹ <https://support.hortonworks.com/s/case/5004400000wHvFbAAK/detail>

²²² <https://support.hortonworks.com/s/detail/0034400001yn8e4AAA>

²²³ <https://support.hortonworks.com/s/case/5004400000wHtF4AAK/detail>

²²⁴ <https://support.hortonworks.com/s/detail/0034400001mobSpAAI>

2	185198 ²²⁵	Hive LLAP Queries are slow	8/20/2018 10:51 PM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Hariprasad Allak ²²⁶	Pablo Idiazquez	<p>simple select count(*) query taking more than 200 seconds time to fetch the result. It executes with high amount of resources which are configured at Yarn Queue manager . Our business expects to reduce the LLAP response time to less than 5 sec for all standard queries as we are depending on Hive LLAP to be our interactive query access mechanism for business users.</p> <p>select count(*) results time improved (Created a new ticket to improve the LLAP performance #00188417)</p>
1	185022 ²²⁷	Performance Issues in the Cluster	8/20/2018 2:20 AM	Closed	Resolved	Hive	S2 - Core Feature Inoperative	Hariprasad Allak ²²⁸	Italo Cocio Figueroa	<p>Fixes:</p> <p>hive-site.xml/custom tez-site.xml.</p> <p>Both places tez.cancel.delegation.tokens.on.completion property is not there in place.</p> <p>With the change in the settings we are retaining this delegation token for a longer time for the hive user to use the same to access hdfs without having any issues.</p>

²²⁵ <https://support.hortonworks.com/s/case/5004400000wHt1RAAS/detail>

²²⁶ <https://support.hortonworks.com/s/detail/0034400001mobSpAAI>

²²⁷ <https://support.hortonworks.com/s/case/5004400000yHRj4AAG/detail>

²²⁸ <https://support.hortonworks.com/s/detail/0034400001mobSpAAI>